



University of Tehran
College of Engineering
School of Electrical and Computer Engineering



Intelligent Systems

Dr.Hosseini

Homework 4

Soroush Mesforush Mashhad

SN:810198472

Dey 01

Contents

1	Question 1 : Multi layered perceptron	5
1.1	Theoretical Analysis	5
1.2	Python Implementation	11
2	Question 2 : Application of neural networks in classification	15
2.1	Prerequisites	15
2.2	Part 1 : Defining the network and testing for different batch sizes	19
2.2.1	The MLP Network	20
2.3	Part 2 : Changing activation functions	24
2.3.1	Sigmoid	24
2.3.2	tanh	27
2.3.3	Selu	30
2.4	Part 3 : Changing the loss functions	34
2.4.1	Poisson loss	34
2.4.2	Mean Squared Error loss	37
2.5	Part 4 : Changing the optimizers	41
2.5.1	Adam	41
2.5.2	RMSprop	42
2.6	Part 5 : The best model	43
2.7	Part 1 (CNN) : The preliminary CNN	46
2.8	Part 2 (CNN) : Pooling and Batch Normalization	47
2.8.1	Batch Normalization	47
2.8.2	Pooling	47
2.9	Part 3 (CNN) : Adding dropout	52
2.9.1	Dropout	52
2.10	Part 4 (CNN) : Early Stopping	57
2.10.1	Early Stopping	57

3	Question 3 : Transfer Learning for EfficientNet network	62
3.1	Part A : Studying EfficientNet	62
3.1.1	Network Architecture	63
3.1.2	Different EfficientNet Architectures	64
3.1.3	Preliminary preprocessing on input image	69
3.1.4	EfficientNet Pros compared to other models	69
3.2	Part B : Network Implementation with transfer learning . . .	71
3.3	Part C : Fixing a common problem	73
3.4	Part 4 : Retraining the model on a new dataset	75

Abstract

In this project we begin with analyzing a multi layered perceptron neural network theoretically, then we verify our calculations via python.

In the second part we design an **MLP** neural network and change various elements in it and observe the accuracy for different states. Then we go on to add a **CNN** to the **MLP** design and observe the effects of pooling, batch normalization etc...

In the final section we first study EfficientNet then perform transfer learning on it.

1 Question 1 : Multi layered perceptron

1.1 Theoretical Analysis

Here we shall update the coefficients for two iterations.




Figure 1: Multi Layered Perceptron

For my *SID* : 810198472 the values shall be as depicted below.

$$X_1 = \begin{bmatrix} 2 \\ 7 \\ 2 \end{bmatrix}, Y_1 = \begin{bmatrix} 2 \end{bmatrix}, \quad X_2 = \begin{bmatrix} 7 \\ 2 \\ 7 \end{bmatrix}, Y_2 = \begin{bmatrix} 7 \end{bmatrix}$$

$$W_1 = \begin{bmatrix} 0.12 & 0.27 \\ 0.32 & 0.47 \\ 0.52 & 0.47 \end{bmatrix}, \quad W_2 = \begin{bmatrix} 2.15 & 2.25 & 2.35 \\ 7.45 & 7.55 & 7.65 \end{bmatrix}$$

$$W_3 = \begin{bmatrix} 14.12 & 14.22 \\ 14.32 & 14.42 \\ 14.52 & 14.62 \end{bmatrix}, \quad W_4 = \begin{bmatrix} -4.84 \\ -4.64 \end{bmatrix}$$

$$B_1 = \begin{bmatrix} 0.21 \\ 0.72 \end{bmatrix}, \quad B_2 = \begin{bmatrix} 9.15 \\ 9.25 \\ 9.35 \end{bmatrix}, \quad B_3 = \begin{bmatrix} 0.37 \\ 0.47 \end{bmatrix}, \quad B_4 = \begin{bmatrix} 5.26 \end{bmatrix}$$

The error function is depicted as below.

$$E = \frac{1}{2}(\hat{y} - y)^2$$

We shall first start the feed forward with with X_1 and Y_1 as depicted below.

$$\begin{aligned} Z &= \tanh(W_1^T X + B_1), \quad W_1^T X + B_1 = \begin{bmatrix} 0.12 & 0.32 & 0.52 \\ 0.27 & 0.47 & 0.47 \end{bmatrix} \begin{bmatrix} 2 \\ 7 \\ 2 \end{bmatrix} + \begin{bmatrix} 0.21 \\ 0.72 \end{bmatrix} \\ &\longrightarrow \begin{bmatrix} 3.52 \\ 4.77 \end{bmatrix} + \begin{bmatrix} 0.21 \\ 0.72 \end{bmatrix} = \begin{bmatrix} 3.73 \\ 5.49 \end{bmatrix}, \quad Z = \tanh \left(\begin{bmatrix} 3.73 \\ 5.49 \end{bmatrix} \right) \longrightarrow Z = \begin{bmatrix} 0.99885 \\ 0.99997 \end{bmatrix} \\ K &= \text{Sigmoid}(W_2^T Z + B_2), \quad W_2^T Z + B_2 = \begin{bmatrix} 2.15 & 7.45 \\ 2.25 & 7.55 \\ 2.35 & 7.65 \end{bmatrix} \begin{bmatrix} 0.99885 \\ 0.99997 \end{bmatrix} + \begin{bmatrix} 9.15 \\ 9.25 \\ 9.35 \end{bmatrix} \\ &\longrightarrow \begin{bmatrix} 9.597 \\ 9.797 \\ 9.997 \end{bmatrix} + \begin{bmatrix} 9.15 \\ 9.25 \\ 9.35 \end{bmatrix} = \begin{bmatrix} 18.747 \\ 19.038 \\ 19.339 \end{bmatrix}, \quad K = \text{Sigmoid} \left(\begin{bmatrix} 18.747 \\ 19.038 \\ 19.339 \end{bmatrix} \right) = \begin{bmatrix} 0.9999 \\ 0.9999 \\ 0.9999 \end{bmatrix} \\ P &= \tanh(W_3^T K + B_3), \quad W_3^T K + B_3 = \begin{bmatrix} 14.12 & 14.32 & 14.52 \\ 14.22 & 14.42 & 14.62 \end{bmatrix} \begin{bmatrix} 0.9999 \\ 0.9999 \\ 0.9999 \end{bmatrix} + \begin{bmatrix} 0.37 \\ 0.47 \end{bmatrix} \\ &\longrightarrow \begin{bmatrix} 42.956 \\ 43.256 \end{bmatrix} + \begin{bmatrix} 0.37 \\ 0.47 \end{bmatrix} = \begin{bmatrix} 43.326 \\ 43.726 \end{bmatrix}, \quad P = \tanh \left(\begin{bmatrix} 43.326 \\ 43.726 \end{bmatrix} \right) = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \\ \hat{y} &= \text{Relu}(W_4^T P + B_4), \quad W_4^T P + B_4 = \begin{bmatrix} -4.84 & -4.64 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} + \begin{bmatrix} 5.26 \end{bmatrix} = \begin{bmatrix} -4.22 \end{bmatrix}. \\ \hat{y} &= \text{Relu}(-4.22) = \max(-4.22, 0) = 0, \quad E = \frac{1}{2}(\hat{y} - y)^2 = 2 \end{aligned}$$

Now we shall perform back propagation. This was very difficult for me and I watched [Andrew Ng's](#) neural network's course. We studied in the lectures that we update parameters as follows.

$$W_n = W_n - lr \frac{\partial L}{\partial W_n}, \quad B_n = B_n - lr \frac{\partial L}{\partial B_n}$$

In the following picture I have included [Andrew Ng's](#) handwritten lecture of the backpropagation equations then I have tidied them up accordingly.

Backward propagation for layer l

→ Input $da^{[l]}$

→ Output $da^{[l-1]}, dW^{[l]}, db^{[l]}$

$$\begin{aligned}
 dz^{[l]} &= da^{[l]} * g'^{[l]}(z^{[l]}) \\
 dW^{[l]} &= dz^{[l]} \cdot a^{[l-1]T} \\
 db^{[l]} &= dz^{[l]} \\
 da^{[l-1]} &= W^{[l]T} \cdot dz^{[l]} \\
 dz^{[l+1]} &= W^{[l+1]T} dz^{[l]} * g'^{[l+1]}(z^{[l+1]})
 \end{aligned}$$

Figure 2: Backpropagation equations

$$\begin{aligned}
 \frac{\partial L}{\partial z^{[l]}} &= \frac{\partial L}{\partial a^{[l]}} \circ \frac{\partial a^{[l]}}{\partial z^{[l]}} \\
 \frac{\partial L}{\partial W^{[l]}} &= \frac{\partial L}{\partial z^{[l]}} \cdot (a^{[l-1]})^T \\
 \frac{\partial L}{\partial B^{[l]}} &= \frac{\partial L}{\partial z^{[l]}} \\
 \frac{\partial L}{\partial a^{[l-1]}} &= W^{[l]T} \frac{\partial L}{\partial z^{[l]}} \\
 z^{[l]} &= W^{[l]} \cdot a^{[l-1]} + b^{[l]}
 \end{aligned}$$

Now we shall continue to find the necessary derivatives, before doing so we

must pay attention that.

$$f(z) = \frac{1}{1 + e^{-z}} \longrightarrow f'(z) = 1 - f(z),$$

$$f(z) = \tanh(z) \longrightarrow f'(z) = 1 - f^2(z)$$

$$f(z) = \text{Relu}(z) \longrightarrow f'(z) = H(z), \quad H(z) : \text{Heaviside}$$

Now we shall perform some very lengthy derivatives.

Layer 4

$$\begin{aligned} \frac{\partial L}{\partial W^{[4]}}{}^T &= \frac{\partial L}{\partial z^{[4]}} \cdot (a^{[3]})^T, \quad \frac{\partial L}{\partial z^{[4]}} = \frac{\partial L}{\partial a^{[4]}} \circ \frac{\partial a^{[4]}}{\partial z^{[4]}} = (\hat{y} - y)H(W_4^T P + B_4) \\ \frac{\partial L}{\partial W^{[4]}}{}^T &= (\hat{y} - y)H(W_4^T P + B_4)P^T, \quad \frac{\partial L}{\partial B^{[4]}} = (\hat{y} - y)H(W_4^T P + B_4) \end{aligned}$$

Layer 3

$$\begin{aligned} \frac{\partial L}{\partial W^{[3]}}{}^T &= \frac{\partial L}{\partial z^{[3]}} \cdot (a^{[2]})^T, \quad \frac{\partial L}{\partial z^{[3]}} = \frac{\partial L}{\partial a^{[3]}} \circ \frac{\partial a^{[3]}}{\partial z^{[3]}} = (W_4(\hat{y} - y)H(W_4^T P + B_4)) \circ (1 - P^2) \\ \frac{\partial L}{\partial W^{[3]}}{}^T &= ((W_4(\hat{y} - y)H(W_4^T P + B_4)) \circ (1 - P^2)) \cdot K^T, \\ \frac{\partial L}{\partial B^{[3]}} &= (W_4(\hat{y} - y)H(W_4^T P + B_4)) \circ (1 - P^2) \end{aligned}$$

Layer 2

$$\begin{aligned} \frac{\partial L}{\partial W^{[2]}}{}^T &= \frac{\partial L}{\partial z^{[2]}} \cdot (a^{[1]})^T, \\ \frac{\partial L}{\partial z^{[2]}} &= \frac{\partial L}{\partial a^{[2]}} \circ \frac{\partial a^{[2]}}{\partial z^{[2]}} = (W_3(W_4(\hat{y} - y)H(W_4^T P + B_4)) \circ (1 - P^2)) \circ (K(1 - K)) \\ \frac{\partial L}{\partial W^{[2]}}{}^T &= (W_3(W_4(\hat{y} - y)H(W_4^T P + B_4)) \circ (1 - P^2)) \circ (K(1 - K)) \cdot Z^T, \\ \frac{\partial L}{\partial B^{[2]}} &= (W_3(W_4(\hat{y} - y)H(W_4^T P + B_4)) \circ (1 - P^2)) \circ (K(1 - K)) \end{aligned}$$

Layer 1

$$\frac{\partial L}{\partial W^{[1]}}{}^T = \frac{\partial L}{\partial z^{[1]}} \cdot (a^{[0]})^T,$$

$$\frac{\partial L}{\partial z^{[1]}} = \frac{\partial L}{\partial a^{[1]}} \circ \frac{\partial a^{[1]}}{\partial z^{[2]}} = (W_2 (W_3 (W_4 (\hat{y} - y) H(W_4^T P + B_4)) \circ (1 - P^2)) \circ (K(1 - K))) \circ (1 - Z^2)$$

$$\frac{\partial L}{\partial W^{[1]}}{}^T = (W_2 (W_3 (W_4 (\hat{y} - y) H(W_4^T P + B_4)) \circ (1 - P^2)) \circ (K(1 - K))) \circ (1 - Z^2) \cdot X^T,$$

$$\frac{\partial L}{\partial B^{[1]}} = (W_2 (W_3 (W_4 (\hat{y} - y) H(W_4^T P + B_4)) \circ (1 - P^2)) \circ (K(1 - K))) \circ (1 - Z^2)$$

Now for the first backpropagation we have:

$$\frac{\partial L}{\partial W^{[4]}} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}, \quad \frac{\partial L}{\partial B^{[4]}} = [0]$$

$$\frac{\partial L}{\partial W^{[3]}} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}, \quad \frac{\partial L}{\partial B^{[3]}} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\frac{\partial L}{\partial W^{[2]}} = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix}, \quad \frac{\partial L}{\partial B^{[2]}} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

$$\frac{\partial L}{\partial W^{[1]}} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}, \quad \frac{\partial L}{\partial B^{[1]}} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

The results are somewhat unusual, I have double and triple checked my answers and there is nothing wrong with it, the results are like this because of my SID, so when we perform the updating, nothing happens, we shall go on to give X_2 as the input.

$$\begin{aligned}
Z &= \tanh(W_1^T X + B_1), \quad W_1^T X + B_1 = \begin{bmatrix} 0.12 & 0.32 & 0.52 \\ 0.27 & 0.47 & 0.47 \end{bmatrix} \begin{bmatrix} 7 \\ 2 \\ 7 \end{bmatrix} + \begin{bmatrix} 0.21 \\ 0.72 \end{bmatrix} \\
&\longrightarrow \begin{bmatrix} 5.12 \\ 6.12 \end{bmatrix} + \begin{bmatrix} 0.21 \\ 0.72 \end{bmatrix} = \begin{bmatrix} 5.33 \\ 6.84 \end{bmatrix}, \quad Z = \tanh \left(\begin{bmatrix} 5.33 \\ 6.84 \end{bmatrix} \right) \longrightarrow Z = \begin{bmatrix} 0.99995 \\ 0.99999 \end{bmatrix} \\
K &= \text{Sigmoid}(W_2^T Z + B_2), \quad W_2^T Z + B_2 = \begin{bmatrix} 2.15 & 7.45 \\ 2.25 & 7.55 \\ 2.35 & 7.65 \end{bmatrix} \begin{bmatrix} 0.99995 \\ 0.99999 \end{bmatrix} + \begin{bmatrix} 9.15 \\ 9.25 \\ 9.35 \end{bmatrix} \\
&\longrightarrow \begin{bmatrix} 9.6 \\ 9.8 \\ 9.99 \end{bmatrix} + \begin{bmatrix} 9.15 \\ 9.25 \\ 9.35 \end{bmatrix} = \begin{bmatrix} 18.75 \\ 19.05 \\ 19.34 \end{bmatrix}, \quad K = \text{Sigmoid} \left(\begin{bmatrix} 18.75 \\ 19.05 \\ 19.34 \end{bmatrix} \right) = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} \\
P &= \tanh(W_3^T K + B_3), \quad W_3^T K + B_3 = \begin{bmatrix} 14.12 & 14.32 & 14.52 \\ 14.22 & 14.42 & 14.62 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} + \begin{bmatrix} 0.37 \\ 0.47 \end{bmatrix} \\
&\longrightarrow \begin{bmatrix} 42.96 \\ 43.26 \end{bmatrix} + \begin{bmatrix} 0.37 \\ 0.47 \end{bmatrix} = \begin{bmatrix} 43.33 \\ 43.73 \end{bmatrix}, \quad P = \tanh \left(\begin{bmatrix} 43.33 \\ 43.73 \end{bmatrix} \right) = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \\
\hat{y} &= \text{Relu}(W_4^T P + B_4), \quad W_4^T P + B_4 = \begin{bmatrix} -4.84 & -4.64 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} + \begin{bmatrix} 5.26 \end{bmatrix} = \begin{bmatrix} -4.22 \end{bmatrix}. \\
\hat{y} &= \text{Relu}(-4.22) = \max(-4.22, 0) = 0, \quad E = \frac{1}{2}(\hat{y} - y)^2 = 24.5
\end{aligned}$$

The backpropagation is exactly like the previous iteration, the vanishing gradient problem happens again.

$$\frac{\partial L}{\partial W^{[4]}} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}, \quad \frac{\partial L}{\partial B^{[4]}} = [0]$$

$$\frac{\partial L}{\partial W^{[3]}} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}, \quad \frac{\partial L}{\partial B^{[3]}} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\frac{\partial L}{\partial W^{[2]}} = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix}, \quad \frac{\partial L}{\partial B^{[2]}} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

$$\frac{\partial L}{\partial W^{[1]}} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}, \quad \frac{\partial L}{\partial B^{[1]}} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

1.2 Python Implementation

Here we have implemented what we did theoretically in python, the codes are as below.

```

1 def step(x):
2     if x>0:
3         return 1
4     else:
5         return 0
6 
```

Figure 3: Step function

```

1 def AclayFeed(x,W1,W2,W3,W4,B1,B2,B3,B4):
2     Z = np.tanh(np.dot(np.transpose(W1),x)+B1)
3     K= 1/(1+np.exp(-(np.dot(np.transpose(W2),Z)+B2)))
4     P = np.tanh(np.dot(np.transpose(W3),K)+B3)
5     Yhat = np.maximum(0,np.dot(np.transpose(W4),P)+B4)
6     return Yhat[0][0]
7 
```

Figure 4: Yhat calculator

```

1  def Gradient_CalcandUp(x,W1,W2,W3,W4,B1,B2,B3,B4,Y,lr):
2      Z = np.tanh(np.dot(np.transpose(W1),x)+B1)
3      K= 1/1+np.exp(-(np.dot(np.transpose(W2),Z)+B2))
4      P = np.tanh(np.dot(np.transpose(W3),K)+B3)
5      Yhat = np.maximum(0,np.dot(np.transpose(W4),P)+B4)
6      A = np.dot(np.transpose(W4),P)+B4
7      dW4 = np.transpose((Yhat - Y)*step(A)*np.transpose(P))
8      dB4 = (Yhat - Y)*step(A)
9      dW3 = np.transpose(np.matmul(np.multiply(W4 *(Yhat - Y)*A,1-P
10         **2),np.transpose(K)))
11      dB3 = np.multiply(W4 *(Yhat - Y)*A,1-P**2)
12      TempW2=np.multiply(W4 *(Yhat - Y)*A,1-P**2)
13      dW2 = np.transpose(np.matmul(np.multiply(np.matmul(W3,TempW2),
14         K*(1-K)),np.transpose(Z)))
15      dB2 =np.multiply(np.matmul(W3,TempW2),K*(1-K))
16      TempW1 = np.multiply(np.matmul(W3,TempW2),K*(1-K))
17      dW1 = np.transpose(np.matmul(np.multiply(np.matmul(W2,TempW1),
18         1-Z**2),np.transpose(x)))
19      dB1 = np.multiply(np.matmul(W2,TempW1),1-Z**2)
20      W4-=lr*dW4
21      W3-=lr*dW3
22      W2-=lr*dW2
23      W1-=lr*dW1
24      B4-=lr*dB4
25      B3-=lr*dB3
26      B2-=lr*dB2
27      B1-=lr*dB1
28      return W4,W3,W2,W1,B4,B3,B2,B1

```

Figure 5: Gradient calculator

After this we shall continue to obtain the final results.

```

1  #First Iter
2  Y1=2
3  Yhat1=AclayFeed(x1,W1,W2,W3,W4,B1,B2,B3,B4)
4  print(f"Error for first feed forward: {ErrorCalc(Y1,Yhat1)}")
5  lr = 0.1
6  W4up1,W3up1,W2up1,W1up1,B4up1,B3up1,B2up1,B1up1=
    Gradient_CalcandUp(x1,W1,W2,W3,W4,B1,B2,B3,B4,Y1,lr)
7  print(f"Updated W4 after 1 iteration : {W4up1}")
8  print(f"Updated W3 after 1 iteration : {W3up1}")
9  print(f"Updated W2 after 1 iteration : {W2up1}")
10 print(f"Updated W1 after 1 iteration : {W1up1}")
11 print(f"Updated B4 after 1 iteration : {B4up1}")
12 print(f"Updated B3 after 1 iteration : {B3up1}")
13 print(f"Updated B2 after 1 iteration : {B2up1}")
14 print(f"Updated B1 after 1 iteration : {B1up1}")
15 ##Second Iter
16 Y2=7
17 Yhat2=AclayFeed(x2,W1,W2,W3,W4,B1,B2,B3,B4)
18 print( '
    _____
    ')
19 print(f"Error for second feed forward: {ErrorCalc(Y2,Yhat2)}")
20 lr = 0.1
21 W4up2,W3up2,W2up2,W1up2,B4up2,B3up2,B2up2,B1up2=
    Gradient_CalcandUp(x2,W1,W2,W3,W4,B1,B2,B3,B4,Y2,lr)
22 print(f"Updated W4 after 2 iterations : {W4up2}")
23 print(f"Updated W3 after 2 iterations : {W3up2}")
24 print(f"Updated W2 after 2 iterations : {W2up2}")
25 print(f"Updated W1 after 2 iterations : {W1up2}")
26 print(f"Updated B4 after 2 iterations : {B4up2}")
27 print(f"Updated B3 after 2 iterations : {B3up2}")
28 print(f"Updated B2 after 2 iterations : {B2up2}")
29 print(f"Updated B1 after 2 iterations : {B1up2}")
30

```

Figure 6: Error calculator

```
1  Error for first feed forward: 2.0
2  Updated W4 after 1 iteration : [[-4.84]
3  [-4.64]]
4  Updated W3 after 1 iteration : [[14.12 14.22]
5  [14.32 14.42]
6  [14.52 14.62]]
7  Updated W2 after 1 iteration : [[2.15 2.25 2.35]
8  [7.45 7.55 7.65]]
9  Updated W1 after 1 iteration : [[0.12 0.27]
10 [0.32 0.47]
11 [0.52 0.47]]
12 Updated B4 after 1 iteration : [[5.26]]
13 Updated B3 after 1 iteration : [[0.37]
14 [0.47]]
15 Updated B2 after 1 iteration : [[9.15]
16 [9.25]
17 [9.35]]
18 Updated B1 after 1 iteration : [[0.21]
19 [0.72]]
20
21 Error for second feed forward: 24.5
22 Updated W4 after 2 iterations : [[-4.84]
23 [-4.64]]
24 Updated W3 after 2 iterations : [[14.12 14.22]
25 [14.32 14.42]
26 [14.52 14.62]]
27 Updated W2 after 2 iterations : [[2.15 2.25 2.35]
28 [7.45 7.55 7.65]]
29 Updated W1 after 2 iterations : [[0.12 0.27]
30 [0.32 0.47]
31 [0.52 0.47]]
32 Updated B4 after 2 iterations : [[5.26]]
33 Updated B3 after 2 iterations : [[0.37]
34 [0.47]]
35 Updated B2 after 2 iterations : [[9.15]
36 [9.25]
37 [9.35]]
38 Updated B1 after 2 iterations : [[0.21]
39 [0.72]]
40
```

Figure 7: Error calculator

2 Question 2 : Application of neural networks in classification

2.1 Prerequisites

First of all, we shall depict the imported things used in this section.

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sn
4 from matplotlib import pyplot as plt
5 import random
6 from matplotlib import pyplot
7 from sklearn.metrics import confusion_matrix
8 from sklearn.metrics import precision_score , recall_score ,
   fl_score , accuracy_score
9 from sklearn.model_selection import train_test_split
10 from keras.datasets import cifar10
11 from keras.callbacks import EarlyStopping
12 from keras.layers import BatchNormalization
13 from tensorflow.keras.utils import to_categorical
14 from tensorflow.keras.models import Sequential
15 from tensorflow.keras.layers import Conv2D , Flatten , Dense ,
   MaxPooling2D , Dropout
16 from tensorflow.keras.optimizers import SGD , Adam
```

Figure 8: Imported libraries

```
1 def Loader():
2     (X_train, Y_train), (X_test, Y_test) = cifar10.load_data()
3     return X_train, Y_train, X_test, Y_test
4
```

Figure 9: Loader function

Now we go on to print the first ten photos of the **Cifar-10** datasets.

```
1  Cifar_Labels = [ 'airplane', 'automobile', 'bird', 'cat', 'deer',  
2                    ', 'dog', 'frog', 'horse', 'ship', 'truck' ]  
3  X_train, Y_train, X_test, Y_test = Loader()  
4  for i in range(10):  
5      plt.subplot(2,5,i+1)  
6      plt.imshow(X_train[i])  
7      plt.title(Cifar_Labels[Y_train[i][0]])
```

Figure 10: First ten photos



Figure 11: First ten photos

Now we shall define the rest of our needed functions including the **MLP**.

```
1  def preprocessing(X_train, Y_train, X_test, Y_test):
2  # Transform label indices to one-hot encoded vectors
3  Y_train = to_categorical(Y_train, num_classes = len(
4      Cifar_Labels))
5  Y_test = to_categorical(Y_test, num_classes = len(Cifar_Labels)
6      )
7  #Because working with 4 dimensions is not possible but is
8  #possible in cnns, we
9  #change the dimensions
10 X_train = np.reshape(X_train, (X_train.shape[0], 32*32*3))
11 X_test = np.reshape(X_test, (X_test.shape[0], 32*32*3))
12 X_train = X_train.astype('float32')
13 X_test = X_test.astype('float32')
14 #Here we perform normalization
15 X_train = X_train/255
16 X_test = X_test/255
17 #Test data must be kept unseen from the model until final
18 #evaluation hence we need
19 #validation dataset.
20 __, X_valid, __, Y_valid = train_test_split(X_train, Y_train,
21     test_size=0.2, random_state= 8)
22 return X_train, Y_train, X_test, Y_test, X_valid, Y_valid
```

Figure 12: Preprocessing

```
1 def Acc_plot(hist):
2     fig = plt.figure()
3     plt.plot(hist.history['accuracy'], 'r')
4     plt.plot(hist.history['val_accuracy'], 'b')
5     plt.title('Model Accuracy')
6     plt.ylabel('Acc')
7     plt.xlabel('Epoch')
8     plt.legend(['Train', 'Validation'])
9     plt.grid()
10
```

Figure 13: Accuracy plot

```
1 def loss_plot(hist):
2     fig = plt.figure()
3     plt.plot(hist.history['loss'], 'r')
4     plt.plot(hist.history['val_loss'], 'b')
5     plt.title('Model Loss')
6     plt.ylabel('Loss')
7     plt.xlabel('Epoch')
8     plt.legend(['Train', 'Validation'])
9     plt.grid()
10
```

Figure 14: Loss plot

After this we have the functions needed to perform analysis after training the models.

```
1 def CompleteAnalysis(model , x_test , y_test ):
2     Y_pred = model.model.predict(x_test)
3     y_pred = np.argmax(Y_pred, axis=1)
4     Y_test = np.argmax(y_test, axis=1)
5     test_loss , test_accuracy = model.model.evaluate(x_test , y_test)
6     print('test loss = %f' % test_loss)
7     print('test accuracy = %f' % test_accuracy)
8     # accuracy: (tp + tn) / (p + n)
9     accuracy = accuracy_score(Y_test, y_pred)
10    print('Accuracy: %f' % accuracy)
11    # precision tp / (tp + fp)
12    precision = precision_score(Y_test, y_pred, average='macro')
13    print('Precision: %f' % precision)
14    # recall: tp / (tp + fn)
15    recall = recall_score(Y_test, y_pred, average='macro')
16    print('Recall: %f' % recall)
17    # f1: 2 tp / (2 tp + fp + fn)
18    f1 = f1_score(Y_test, y_pred, average='macro')
19    print('F1 score: %f' % f1)
20    #confMat
21    cm= confusion_matrix(Y_test, y_pred)
22    fig , ax = plt.subplots(figsize=(10,10))
23    sn.heatmap(cm, annot=True, square = True, ax=ax)
24    plt.xlabel('predicted value')
25    plt.ylabel('true value');
26
```

Figure 15: Complete Analysis(including precision, recall and f1-score)

2.2 Part 1 : Defining the network and testing for different batch sizes

Here we define the neural network as needed and observe the results with three different batch sizes.

2.2.1 The MLP Network

The network is depicted as follows.

```
1 def NeuralNetwork(Ac1,Ac2,Ac3,l,opt):
2     model = Sequential()
3     model.add(Dense(128,activation=Ac1,input_dim = 32*32*3))
4     model.add(Dense(128,activation=Ac2))
5     model.add(Dense(10,activation=Ac3))
6     model.compile(loss=l, optimizer=opt, metrics=['accuracy'])
7     return model
8
```

Figure 16: The MLP Network

In the next few pages I have included the code to train (for one example because they are exactly the same) and the results accordingly.

Batch Size = 32

The code to train is as follows.

```
1 X_train, Y_train, X_test, Y_test = Loader()
2 X_train, Y_train, X_test, Y_test, X_valid, Y_valid =
   preprocessing(X_train, Y_train, X_test, Y_test)
3 model32 = NeuralNetwork('relu', 'relu', 'softmax', '
   categorical_crossentropy', 'sgd')
4 history32 = model32.fit(X_train, Y_train, epochs=15, batch_size
   =32, validation_data=(X_valid, Y_valid), verbose=1)
5
```

Figure 17: Code to train

```

1 Epoch 1/15
2 1563/1563 [=====] - 16s 9ms/step - loss: 1.8948 - accuracy: 0.3189 - val_loss:
   1.7492 - val_accuracy: 0.3777
3 Epoch 2/15
4 1563/1563 [=====] - 15s 9ms/step - loss: 1.7112 - accuracy: 0.3914 - val_loss:
   1.6516 - val_accuracy: 0.4172
5 Epoch 3/15
6 1563/1563 [=====] - 14s 9ms/step - loss: 1.6311 - accuracy: 0.4238 - val_loss:
   1.6506 - val_accuracy: 0.4101
7 Epoch 4/15
8 1563/1563 [=====] - 13s 9ms/step - loss: 1.5785 - accuracy: 0.4421 - val_loss:
   1.7122 - val_accuracy: 0.3937
9 Epoch 5/15
10 1563/1563 [=====] - 13s 8ms/step - loss: 1.5369 - accuracy: 0.4561 - val_loss:
   1.5977 - val_accuracy: 0.4156
11 Epoch 6/15
12 1563/1563 [=====] - 13s 8ms/step - loss: 1.5002 - accuracy: 0.4706 - val_loss:
   1.4773 - val_accuracy: 0.4691
13 Epoch 7/15
14 1563/1563 [=====] - 13s 8ms/step - loss: 1.4681 - accuracy: 0.4800 - val_loss:
   1.4910 - val_accuracy: 0.4694
15 Epoch 8/15
16 1563/1563 [=====] - 11s 7ms/step - loss: 1.4442 - accuracy: 0.4893 - val_loss:
   1.4178 - val_accuracy: 0.4942
17 Epoch 9/15
18 1563/1563 [=====] - 11s 7ms/step - loss: 1.4192 - accuracy: 0.4977 - val_loss:
   1.3889 - val_accuracy: 0.5111
19 Epoch 10/15
20 1563/1563 [=====] - 11s 7ms/step - loss: 1.3985 - accuracy: 0.5057 - val_loss:
   1.4229 - val_accuracy: 0.4960
21 Epoch 11/15
22 1563/1563 [=====] - 10s 7ms/step - loss: 1.3774 - accuracy: 0.5134 - val_loss:
   1.3623 - val_accuracy: 0.5166
23 Epoch 12/15
24 1563/1563 [=====] - 14s 9ms/step - loss: 1.3579 - accuracy: 0.5213 - val_loss:
   1.3133 - val_accuracy: 0.5336
25 Epoch 13/15
26 1563/1563 [=====] - 14s 9ms/step - loss: 1.3422 - accuracy: 0.5243 - val_loss:
   1.3074 - val_accuracy: 0.5373
27 Epoch 14/15
28 1563/1563 [=====] - 14s 9ms/step - loss: 1.3265 - accuracy: 0.5292 - val_loss:
   1.3111 - val_accuracy: 0.5342
29 Epoch 15/15
30 1563/1563 [=====] - 14s 9ms/step - loss: 1.3094 - accuracy: 0.5362 - val_loss:
   1.3515 - val_accuracy: 0.5133
31
32

```

Figure 18: The Results for Batch = 32

Batch Size = 64

```

1 Epoch 1/15
2 782/782 [=====] - 8s 9ms/step - loss: 1.9581 - accuracy: 0.2999 - val_loss:
   1.8961 - val_accuracy: 0.3291
3 Epoch 2/15
4 782/782 [=====] - 8s 10ms/step - loss: 1.7871 - accuracy: 0.3704 - val_loss:
   1.8790 - val_accuracy: 0.3327
5 Epoch 3/15
6 782/782 [=====] - 7s 9ms/step - loss: 1.7086 - accuracy: 0.3974 - val_loss:
   1.8043 - val_accuracy: 0.3522
7 Epoch 4/15
8 782/782 [=====] - 8s 10ms/step - loss: 1.6525 - accuracy: 0.4188 - val_loss:
   1.7471 - val_accuracy: 0.3689
9 Epoch 5/15
10 782/782 [=====] - 7s 9ms/step - loss: 1.6090 - accuracy: 0.4343 - val_loss:
   1.6453 - val_accuracy: 0.4077
11 Epoch 6/15
12 782/782 [=====] - 8s 10ms/step - loss: 1.5742 - accuracy: 0.4462 - val_loss:
   1.6363 - val_accuracy: 0.4111
13 Epoch 7/15
14 782/782 [=====] - 7s 9ms/step - loss: 1.5447 - accuracy: 0.4531 - val_loss:
   1.5514 - val_accuracy: 0.4416
15 Epoch 8/15
16 782/782 [=====] - 7s 10ms/step - loss: 1.5163 - accuracy: 0.4668 - val_loss:
   1.5987 - val_accuracy: 0.4280
17 Epoch 9/15
18 782/782 [=====] - 8s 10ms/step - loss: 1.4940 - accuracy: 0.4732 - val_loss:
   1.5079 - val_accuracy: 0.4683
19 Epoch 10/15
20 782/782 [=====] - 6s 8ms/step - loss: 1.4699 - accuracy: 0.4821 - val_loss:
   1.4759 - val_accuracy: 0.4766
21 Epoch 11/15
22 782/782 [=====] - 7s 9ms/step - loss: 1.4542 - accuracy: 0.4876 - val_loss:
   1.4351 - val_accuracy: 0.4933
23 Epoch 12/15
24 782/782 [=====] - 7s 9ms/step - loss: 1.4374 - accuracy: 0.4934 - val_loss:
   1.5039 - val_accuracy: 0.4692
25 Epoch 13/15
26 782/782 [=====] - 7s 8ms/step - loss: 1.4216 - accuracy: 0.4978 - val_loss:
   1.5382 - val_accuracy: 0.4617
27 Epoch 14/15
28 782/782 [=====] - 7s 8ms/step - loss: 1.4086 - accuracy: 0.5022 - val_loss:
   1.5301 - val_accuracy: 0.4569
29 Epoch 15/15
30 782/782 [=====] - 6s 8ms/step - loss: 1.3912 - accuracy: 0.5099 - val_loss:
   1.6270 - val_accuracy: 0.4274
31
32

```

Figure 19: The Results for Batch = 64

Batch Size = 256

```

1 Epoch 1/15
2 196/196 [=====] - 4s 15ms/step - loss: 2.0790 - accuracy: 0.2523 - val_loss:
   1.9610 - val_accuracy: 0.2967
3 Epoch 2/15
4 196/196 [=====] - 3s 14ms/step - loss: 1.9140 - accuracy: 0.3202 - val_loss:
   1.9085 - val_accuracy: 0.3077
5 Epoch 3/15
6 196/196 [=====] - 3s 13ms/step - loss: 1.8535 - accuracy: 0.3457 - val_loss:
   1.8299 - val_accuracy: 0.3522
7 Epoch 4/15
8 196/196 [=====] - 3s 13ms/step - loss: 1.8113 - accuracy: 0.3626 - val_loss:
   1.8077 - val_accuracy: 0.3552
9 Epoch 5/15
10 196/196 [=====] - 3s 13ms/step - loss: 1.7784 - accuracy: 0.3735 - val_loss:
    1.7590 - val_accuracy: 0.3814
11 Epoch 6/15
12 196/196 [=====] - 3s 14ms/step - loss: 1.7512 - accuracy: 0.3866 - val_loss:
    1.7479 - val_accuracy: 0.3850
13 Epoch 7/15
14 196/196 [=====] - 3s 14ms/step - loss: 1.7260 - accuracy: 0.3956 - val_loss:
    1.8168 - val_accuracy: 0.3464
15 Epoch 8/15
16 196/196 [=====] - 3s 14ms/step - loss: 1.7052 - accuracy: 0.4023 - val_loss:
    1.7585 - val_accuracy: 0.3658
17 Epoch 9/15
18 196/196 [=====] - 3s 13ms/step - loss: 1.6877 - accuracy: 0.4087 - val_loss:
    1.7107 - val_accuracy: 0.3915
19 Epoch 10/15
20 196/196 [=====] - 3s 13ms/step - loss: 1.6681 - accuracy: 0.4144 - val_loss:
    1.6977 - val_accuracy: 0.4014
21 Epoch 11/15
22 196/196 [=====] - 3s 14ms/step - loss: 1.6545 - accuracy: 0.4203 - val_loss:
    1.6607 - val_accuracy: 0.4142
23 Epoch 12/15
24 196/196 [=====] - 3s 14ms/step - loss: 1.6406 - accuracy: 0.4251 - val_loss:
    1.6391 - val_accuracy: 0.4260
25 Epoch 13/15
26 196/196 [=====] - 3s 13ms/step - loss: 1.6238 - accuracy: 0.4327 - val_loss:
    1.6755 - val_accuracy: 0.4127
27 Epoch 14/15
28 196/196 [=====] - 3s 13ms/step - loss: 1.6159 - accuracy: 0.4350 - val_loss:
    1.6197 - val_accuracy: 0.4300
29 Epoch 15/15
30 196/196 [=====] - 3s 14ms/step - loss: 1.6018 - accuracy: 0.4403 - val_loss:
    1.6416 - val_accuracy: 0.4193
31
32

```

Figure 20: The Results for Batch = 256

As we can see the train and validation accuracy decreases with increasing the batch size, larger batch sizes take less time to train and but yield to lesser accuracies.

2.3 Part 2 : Changing activation functions

Here we shall change the activation functions for the layers, 3 times.

2.3.1 Sigmoid

```

1 Epoch 1/15
2 782/782 [=====] - 7s 8ms/step - loss: 2.2887 - accuracy: 0.1620 - val_loss:
   2.2634 - val_accuracy: 0.1799
3 Epoch 2/15
4 782/782 [=====] - 3s 4ms/step - loss: 2.2374 - accuracy: 0.2293 - val_loss:
   2.2067 - val_accuracy: 0.2756
5 Epoch 3/15
6 782/782 [=====] - 3s 4ms/step - loss: 2.1701 - accuracy: 0.2582 - val_loss:
   2.1320 - val_accuracy: 0.2792
7 Epoch 4/15
8 782/782 [=====] - 3s 4ms/step - loss: 2.0983 - accuracy: 0.2745 - val_loss:
   2.0671 - val_accuracy: 0.2806
9 Epoch 5/15
10 782/782 [=====] - 3s 4ms/step - loss: 2.0415 - accuracy: 0.2854 - val_loss:
   2.0198 - val_accuracy: 0.2892
11 Epoch 6/15
12 782/782 [=====] - 3s 4ms/step - loss: 1.9987 - accuracy: 0.2954 - val_loss:
   1.9852 - val_accuracy: 0.3018
13 Epoch 7/15
14 782/782 [=====] - 3s 4ms/step - loss: 1.9657 - accuracy: 0.3025 - val_loss:
   1.9520 - val_accuracy: 0.3062
15 Epoch 8/15
16 782/782 [=====] - 3s 4ms/step - loss: 1.9396 - accuracy: 0.3133 - val_loss:
   1.9298 - val_accuracy: 0.3068
17 Epoch 9/15
18 782/782 [=====] - 3s 3ms/step - loss: 1.9187 - accuracy: 0.3194 - val_loss:
   1.9121 - val_accuracy: 0.3175
19 Epoch 10/15
20 782/782 [=====] - 3s 4ms/step - loss: 1.9022 - accuracy: 0.3254 - val_loss:
   1.8946 - val_accuracy: 0.3250
21 Epoch 11/15
22 782/782 [=====] - 3s 4ms/step - loss: 1.8880 - accuracy: 0.3324 - val_loss:
   1.8808 - val_accuracy: 0.3335
23 Epoch 12/15
24 782/782 [=====] - 3s 3ms/step - loss: 1.8756 - accuracy: 0.3358 - val_loss:
   1.8678 - val_accuracy: 0.3335
25 Epoch 13/15
26 782/782 [=====] - 3s 3ms/step - loss: 1.8640 - accuracy: 0.3407 - val_loss:
   1.8618 - val_accuracy: 0.3393
27 Epoch 14/15
28 782/782 [=====] - 3s 4ms/step - loss: 1.8532 - accuracy: 0.3456 - val_loss:
   1.8479 - val_accuracy: 0.3459
29 Epoch 15/15
30 782/782 [=====] - 3s 4ms/step - loss: 1.8432 - accuracy: 0.3475 - val_loss:
   1.8393 - val_accuracy: 0.3512
31
32

```

Figure 21: Sigmoid Activation

Accuracy Plot




Figure 22: Accuracy Plot

Loss Plot




Figure 23: Loss Plot

Confusion matrix and accuracy and error

```

1 1563/1563 [=====] - 3s 2ms/step
2 1563/1563 [=====] - 3s 2ms/step -
   loss: 1.8407 - accuracy: 0.3526
3 test loss = 1.840652
4 test accuracy = 0.352620
5

```

Figure 24: Accuracy and error for Sigmoid

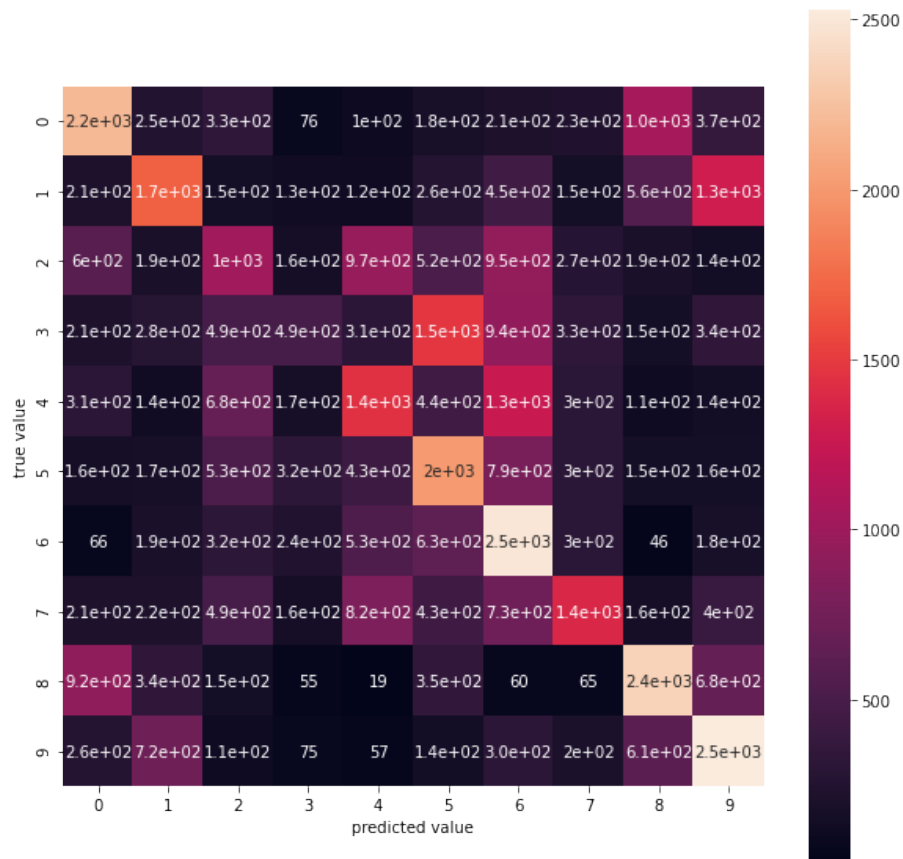


Figure 25: Confusion Matrix

2.3.2 tanh

```

1 Epoch 1/15
2 782/782 [=====] - 3s 4ms/step - loss: 1.9343 - accuracy: 0.3067 - val_loss:
   1.8745 - val_accuracy: 0.3464
3 Epoch 2/15
4 782/782 [=====] - 3s 3ms/step - loss: 1.7746 - accuracy: 0.3737 - val_loss:
   1.7369 - val_accuracy: 0.3925
5 Epoch 3/15
6 782/782 [=====] - 3s 4ms/step - loss: 1.7095 - accuracy: 0.3993 - val_loss:
   1.7124 - val_accuracy: 0.3949
7 Epoch 4/15
8 782/782 [=====] - 3s 4ms/step - loss: 1.6585 - accuracy: 0.4150 - val_loss:
   1.6400 - val_accuracy: 0.4169
9 Epoch 5/15
10 782/782 [=====] - 3s 4ms/step - loss: 1.6211 - accuracy: 0.4273 - val_loss:
    1.6150 - val_accuracy: 0.4242
11 Epoch 6/15
12 782/782 [=====] - 3s 4ms/step - loss: 1.5914 - accuracy: 0.4384 - val_loss:
    1.6131 - val_accuracy: 0.4212
13 Epoch 7/15
14 782/782 [=====] - 3s 4ms/step - loss: 1.5663 - accuracy: 0.4466 - val_loss:
    1.5970 - val_accuracy: 0.4265
15 Epoch 8/15
16 782/782 [=====] - 3s 4ms/step - loss: 1.5446 - accuracy: 0.4541 - val_loss:
    1.5479 - val_accuracy: 0.4451
17 Epoch 9/15
18 782/782 [=====] - 3s 3ms/step - loss: 1.5247 - accuracy: 0.4612 - val_loss:
    1.5227 - val_accuracy: 0.4623
19 Epoch 10/15
20 782/782 [=====] - 3s 4ms/step - loss: 1.5083 - accuracy: 0.4643 - val_loss:
    1.5500 - val_accuracy: 0.4513
21 Epoch 11/15
22 782/782 [=====] - 3s 3ms/step - loss: 1.4910 - accuracy: 0.4764 - val_loss:
    1.6814 - val_accuracy: 0.3894
23 Epoch 12/15
24 782/782 [=====] - 3s 3ms/step - loss: 1.4766 - accuracy: 0.4782 - val_loss:
    1.5115 - val_accuracy: 0.4608
25 Epoch 13/15
26 782/782 [=====] - 3s 3ms/step - loss: 1.4613 - accuracy: 0.4850 - val_loss:
    1.5374 - val_accuracy: 0.4525
27 Epoch 14/15
28 782/782 [=====] - 3s 4ms/step - loss: 1.4452 - accuracy: 0.4891 - val_loss:
    1.4481 - val_accuracy: 0.4845
29 Epoch 15/15
30 782/782 [=====] - 3s 3ms/step - loss: 1.4311 - accuracy: 0.4957 - val_loss:
    1.5060 - val_accuracy: 0.4570
31
32

```

Figure 26: tanh Activation

Accuracy Plot




Figure 27: Accuracy Plot

Loss Plot




Figure 28: Loss Plot

Confusion matrix and accuracy and error

```

1 1563/1563 [=====] - 2s 2ms/step
2 1563/1563 [=====] - 2s 2ms/step -
   loss: 1.5082 - accuracy: 0.4604
3 test loss = 1.508156
4 test accuracy = 0.460420
5

```

Figure 29: Accuracy and error for tanh

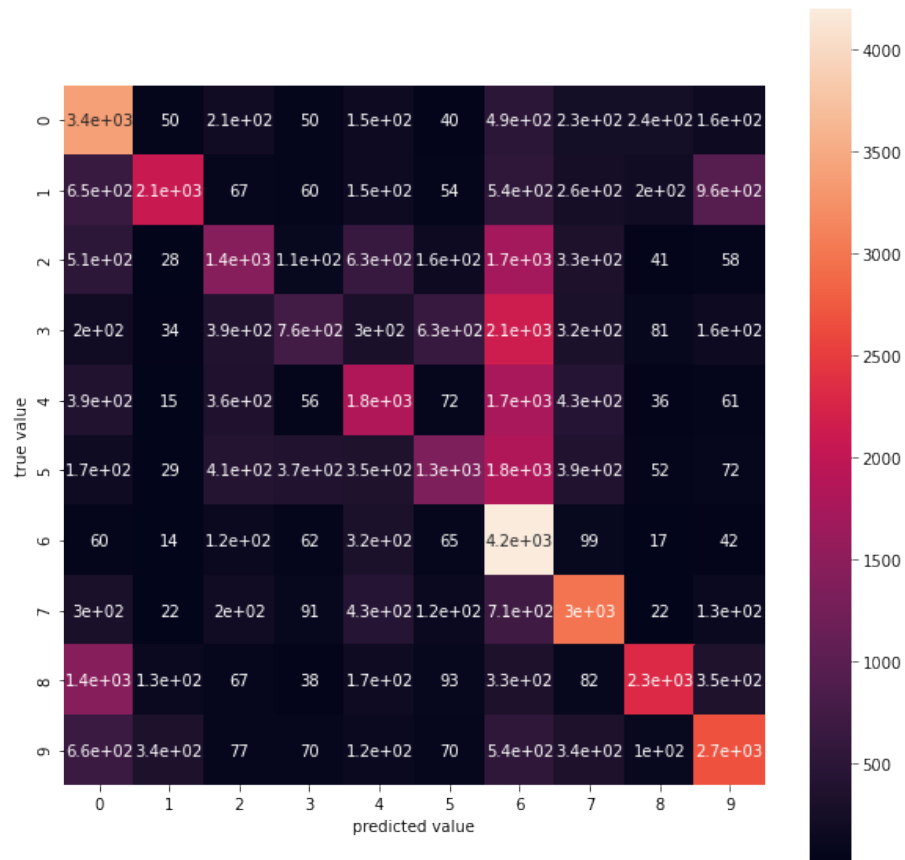


Figure 30: Confusion Matrix

2.3.3 Selu

```

1 Epoch 1/15
2 782/782 [=====] - 3s 4ms/step - loss: 1.9480 - accuracy: 0.3061 - val_loss:
   1.9291 - val_accuracy: 0.2995
3 Epoch 2/15
4 782/782 [=====] - 3s 4ms/step - loss: 1.7708 - accuracy: 0.3736 - val_loss:
   1.9125 - val_accuracy: 0.3250
5 Epoch 3/15
6 782/782 [=====] - 3s 4ms/step - loss: 1.7001 - accuracy: 0.4004 - val_loss:
   1.6834 - val_accuracy: 0.4030
7 Epoch 4/15
8 782/782 [=====] - 3s 4ms/step - loss: 1.6496 - accuracy: 0.4167 - val_loss:
   1.8410 - val_accuracy: 0.3571
9 Epoch 5/15
10 782/782 [=====] - 3s 4ms/step - loss: 1.6107 - accuracy: 0.4302 - val_loss:
   1.6446 - val_accuracy: 0.4066
11 Epoch 6/15
12 782/782 [=====] - 3s 4ms/step - loss: 1.5779 - accuracy: 0.4443 - val_loss:
   1.6786 - val_accuracy: 0.4084
13 Epoch 7/15
14 782/782 [=====] - 3s 4ms/step - loss: 1.5564 - accuracy: 0.4507 - val_loss:
   1.6879 - val_accuracy: 0.3967
15 Epoch 8/15
16 782/782 [=====] - 3s 4ms/step - loss: 1.5307 - accuracy: 0.4598 - val_loss:
   1.8416 - val_accuracy: 0.3590
17 Epoch 9/15
18 782/782 [=====] - 3s 3ms/step - loss: 1.5115 - accuracy: 0.4680 - val_loss:
   1.5949 - val_accuracy: 0.4115
19 Epoch 10/15
20 782/782 [=====] - 3s 4ms/step - loss: 1.4937 - accuracy: 0.4730 - val_loss:
   1.5157 - val_accuracy: 0.4605
21 Epoch 11/15
22 782/782 [=====] - 3s 4ms/step - loss: 1.4725 - accuracy: 0.4814 - val_loss:
   1.5171 - val_accuracy: 0.4559
23 Epoch 12/15
24 782/782 [=====] - 3s 4ms/step - loss: 1.4564 - accuracy: 0.4866 - val_loss:
   1.5869 - val_accuracy: 0.4358
25 Epoch 13/15
26 782/782 [=====] - 3s 4ms/step - loss: 1.4400 - accuracy: 0.4924 - val_loss:
   1.4580 - val_accuracy: 0.4755
27 Epoch 14/15
28 782/782 [=====] - 3s 4ms/step - loss: 1.4286 - accuracy: 0.4980 - val_loss:
   1.4395 - val_accuracy: 0.4798
29 Epoch 15/15
30 782/782 [=====] - 3s 4ms/step - loss: 1.4134 - accuracy: 0.5022 - val_loss:
   1.4289 - val_accuracy: 0.4892
31
32

```

Figure 31: Selu Activation

Accuracy Plot




Figure 32: Accuracy Plot

Loss Plot




Figure 33: Loss Plot

Confusion matrix and accuracy and error

```

1 1563/1563 [=====] - 2s 2ms/step
2 1563/1563 [=====] - 3s 2ms/step -
   loss: 1.4360 - accuracy: 0.4880
3 test loss = 1.436009
4 test accuracy = 0.488000
5

```

Figure 34: Accuracy and error for Selu

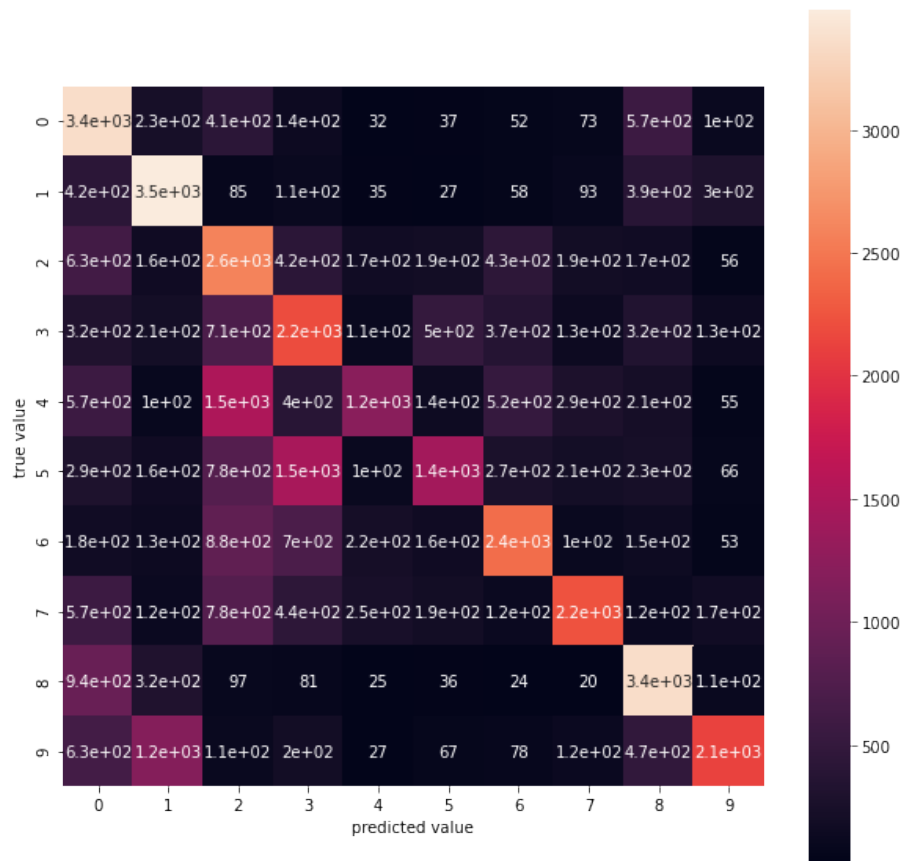


Figure 35: Confusion Matrix

As we can see still the *relu* activation works better than the rest. Now we shall talk about the pros and cons of the used activations here.

Sigmoid Advantages

- It is one of the best normalized functions
- Has a clear prediction
- The gradient is smooth

Sigmoid Disadvantages

- Computationally expensive
- Vanishing gradient can occur
- Always gives positive values

tanh Advantages

- It is zero centric
- The gradient is smooth

tanh Disadvantages

- Computationally expensive
- Vanishing gradient can occur

Selu

It is very similar to Relu but with an advantage, internal normalization is faster than external normalization, which means the network converges faster.

2.4 Part 3 : Changing the loss functions

Here we change the error function, 2 times.

2.4.1 Poisson loss

```

1 Epoch 1/15
2 782/782 [=====] - 4s 4ms/step - loss: 0.3218 - accuracy: 0.2045 - val_loss:
   0.3132 - val_accuracy: 0.2530
3 Epoch 2/15
4 782/782 [=====] - 3s 3ms/step - loss: 0.3071 - accuracy: 0.2743 - val_loss:
   0.3025 - val_accuracy: 0.2835
5 Epoch 3/15
6 782/782 [=====] - 3s 3ms/step - loss: 0.2987 - accuracy: 0.3042 - val_loss:
   0.2958 - val_accuracy: 0.3052
7 Epoch 4/15
8 782/782 [=====] - 3s 3ms/step - loss: 0.2935 - accuracy: 0.3208 - val_loss:
   0.2919 - val_accuracy: 0.3212
9 Epoch 5/15
10 782/782 [=====] - 3s 3ms/step - loss: 0.2900 - accuracy: 0.3358 - val_loss:
    0.2886 - val_accuracy: 0.3381
11 Epoch 6/15
12 782/782 [=====] - 3s 3ms/step - loss: 0.2872 - accuracy: 0.3466 - val_loss:
    0.2861 - val_accuracy: 0.3450
13 Epoch 7/15
14 782/782 [=====] - 3s 3ms/step - loss: 0.2850 - accuracy: 0.3554 - val_loss:
    0.2841 - val_accuracy: 0.3541
15 Epoch 8/15
16 782/782 [=====] - 3s 3ms/step - loss: 0.2831 - accuracy: 0.3610 - val_loss:
    0.2824 - val_accuracy: 0.3549
17 Epoch 9/15
18 782/782 [=====] - 3s 3ms/step - loss: 0.2814 - accuracy: 0.3671 - val_loss:
    0.2809 - val_accuracy: 0.3637
19 Epoch 10/15
20 782/782 [=====] - 3s 3ms/step - loss: 0.2800 - accuracy: 0.3717 - val_loss:
    0.2796 - val_accuracy: 0.3657
21 Epoch 11/15
22 782/782 [=====] - 3s 3ms/step - loss: 0.2786 - accuracy: 0.3753 - val_loss:
    0.2781 - val_accuracy: 0.3746
23 Epoch 12/15
24 782/782 [=====] - 2s 3ms/step - loss: 0.2774 - accuracy: 0.3807 - val_loss:
    0.2770 - val_accuracy: 0.3783
25 Epoch 13/15
26 782/782 [=====] - 2s 3ms/step - loss: 0.2763 - accuracy: 0.3836 - val_loss:
    0.2757 - val_accuracy: 0.3815
27 Epoch 14/15
28 782/782 [=====] - 2s 3ms/step - loss: 0.2752 - accuracy: 0.3884 - val_loss:
    0.2753 - val_accuracy: 0.3846
29 Epoch 15/15
30 782/782 [=====] - 3s 3ms/step - loss: 0.2742 - accuracy: 0.3930 - val_loss:
    0.2734 - val_accuracy: 0.3929
31
32

```

Figure 36: Poissin Loss

Accuracy Plot




Figure 37: Accuracy Plot

Loss Plot




Figure 38: Loss Plot

Confusion matrix and accuracy and error

```

1 1563/1563 [=====] - 2s 1ms/step
2 1563/1563 [=====] - 3s 2ms/step -
   loss: 0.2735 - accuracy: 0.3956
3 test loss = 0.273509
4 test accuracy = 0.395600
5

```

Figure 39: Accuracy and error for Poisson Loss

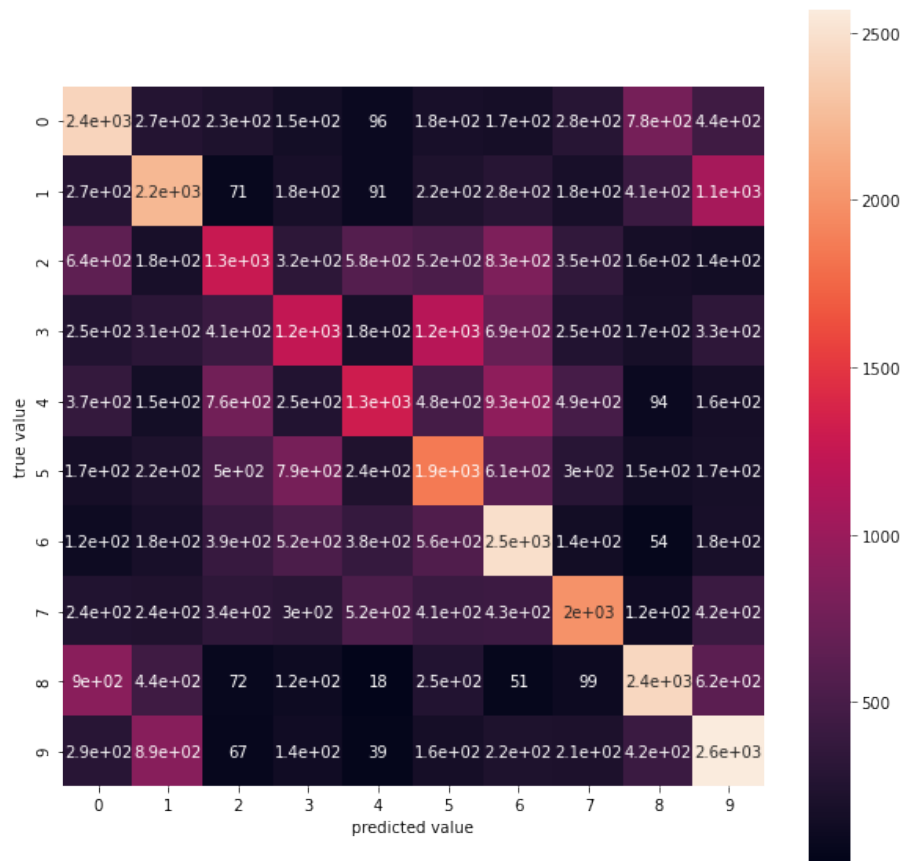


Figure 40: Confusion Matrix

2.4.2 Mean Squared Error loss

```

1 Epoch 1/15
2 782/782 [=====] - 3s 4ms/step - loss: 0.0899 - accuracy: 0.1273 - val_loss:
   0.0892 - val_accuracy: 0.1582
3 Epoch 2/15
4 782/782 [=====] - 3s 4ms/step - loss: 0.0886 - accuracy: 0.1836 - val_loss:
   0.0880 - val_accuracy: 0.2028
5 Epoch 3/15
6 782/782 [=====] - 3s 4ms/step - loss: 0.0876 - accuracy: 0.2138 - val_loss:
   0.0871 - val_accuracy: 0.2289
7 Epoch 4/15
8 782/782 [=====] - 3s 4ms/step - loss: 0.0866 - accuracy: 0.2361 - val_loss:
   0.0862 - val_accuracy: 0.2438
9 Epoch 5/15
10 782/782 [=====] - 3s 3ms/step - loss: 0.0858 - accuracy: 0.2477 - val_loss:
    0.0854 - val_accuracy: 0.2546
11 Epoch 6/15
12 782/782 [=====] - 3s 3ms/step - loss: 0.0850 - accuracy: 0.2574 - val_loss:
    0.0847 - val_accuracy: 0.2590
13 Epoch 7/15
14 782/782 [=====] - 3s 3ms/step - loss: 0.0843 - accuracy: 0.2663 - val_loss:
    0.0840 - val_accuracy: 0.2690
15 Epoch 8/15
16 782/782 [=====] - 3s 3ms/step - loss: 0.0837 - accuracy: 0.2763 - val_loss:
    0.0834 - val_accuracy: 0.2780
17 Epoch 9/15
18 782/782 [=====] - 3s 3ms/step - loss: 0.0831 - accuracy: 0.2859 - val_loss:
    0.0828 - val_accuracy: 0.2847
19 Epoch 10/15
20 782/782 [=====] - 2s 3ms/step - loss: 0.0825 - accuracy: 0.2931 - val_loss:
    0.0823 - val_accuracy: 0.2980
21 Epoch 11/15
22 782/782 [=====] - 2s 3ms/step - loss: 0.0820 - accuracy: 0.3019 - val_loss:
    0.0819 - val_accuracy: 0.3028
23 Epoch 12/15
24 782/782 [=====] - 2s 3ms/step - loss: 0.0816 - accuracy: 0.3096 - val_loss:
    0.0814 - val_accuracy: 0.3098
25 Epoch 13/15
26 782/782 [=====] - 2s 3ms/step - loss: 0.0812 - accuracy: 0.3174 - val_loss:
    0.0810 - val_accuracy: 0.3153
27 Epoch 14/15
28 782/782 [=====] - 2s 3ms/step - loss: 0.0808 - accuracy: 0.3239 - val_loss:
    0.0807 - val_accuracy: 0.3218
29 Epoch 15/15
30 782/782 [=====] - 2s 3ms/step - loss: 0.0804 - accuracy: 0.3283 - val_loss:
    0.0802 - val_accuracy: 0.3269
31
32

```

Figure 41: Mean Squared Error Loss

Accuracy Plot




Figure 42: Accuracy Plot

Loss Plot




Figure 43: Loss Plot

Confusion matrix and accuracy and error

```

1 1563/1563 [=====] - 2s 1ms/step
2 1563/1563 [=====] - 3s 2ms/step -
   loss: 0.0802 - accuracy: 0.3327
3 test loss = 0.080179
4 test accuracy = 0.332700
5

```

Figure 44: Accuracy and error for Mean Squared Error Loss



Figure 45: Confusion Matrix

Our results with `categorical_crossentropy` is still much better than these other two loss functions, I guess this because of the nature of our problem and the dataset we are using, for this problem the `categorical_crossentropy` loss is much better due to us using the **one-hot encoded** version of labels.

Poisson Regression

The Poisson regression is best used for modeling events where the outcomes are counts. To be more precise this loss function is best used when the outcomes are count data, something like non-negative integers, in this problem this is not our case therefore we don't get very good results.

Mean Squared Loss

As we know and can also interpret from the name of this loss function, it is normally used in regression task where the goal is to minimize the expected value of some sort of function on our training data, as shown in the **Stochastic Processes** course this loss is formulated as below.

$$MSE = \frac{1}{n} \sum (y - \hat{y})^2$$

In other words this loss function is used when we want to check how close are estimates are to actual values.

In our project, comparing all loss functions, the best one shall be the **`categorical_crossentropy`** loss.

2.5 Part 4 : Changing the optimizers

Here we shall change the optimizer of the network, 2 times.

2.5.1 Adam

```

1 Epoch 1/15
2 782/782 [=====] - 5s 6ms/step - loss: 1.8580 - accuracy: 0.3291 - val_loss:
   1.7640 - val_accuracy: 0.3550
3 Epoch 2/15
4 782/782 [=====] - 4s 5ms/step - loss: 1.6807 - accuracy: 0.3988 - val_loss:
   1.6233 - val_accuracy: 0.4169
5 Epoch 3/15
6 782/782 [=====] - 4s 5ms/step - loss: 1.6076 - accuracy: 0.4239 - val_loss:
   1.5285 - val_accuracy: 0.4551
7 Epoch 4/15
8 782/782 [=====] - 4s 5ms/step - loss: 1.5559 - accuracy: 0.4450 - val_loss:
   1.5821 - val_accuracy: 0.4325
9 Epoch 5/15
10 782/782 [=====] - 4s 5ms/step - loss: 1.5287 - accuracy: 0.4573 - val_loss:
   1.4826 - val_accuracy: 0.4731
11 Epoch 6/15
12 782/782 [=====] - 4s 5ms/step - loss: 1.4997 - accuracy: 0.4656 - val_loss:
   1.4845 - val_accuracy: 0.4711
13 Epoch 7/15
14 782/782 [=====] - 4s 6ms/step - loss: 1.4774 - accuracy: 0.4738 - val_loss:
   1.4901 - val_accuracy: 0.4686
15 Epoch 8/15
16 782/782 [=====] - 4s 6ms/step - loss: 1.4579 - accuracy: 0.4801 - val_loss:
   1.4195 - val_accuracy: 0.4884
17 Epoch 9/15
18 782/782 [=====] - 4s 5ms/step - loss: 1.4302 - accuracy: 0.4893 - val_loss:
   1.4157 - val_accuracy: 0.4881
19 Epoch 10/15
20 782/782 [=====] - 4s 5ms/step - loss: 1.4206 - accuracy: 0.4903 - val_loss:
   1.4093 - val_accuracy: 0.4929
21 Epoch 11/15
22 782/782 [=====] - 4s 6ms/step - loss: 1.3992 - accuracy: 0.5002 - val_loss:
   1.3666 - val_accuracy: 0.5150
23 Epoch 12/15
24 782/782 [=====] - 4s 5ms/step - loss: 1.3915 - accuracy: 0.4996 - val_loss:
   1.3928 - val_accuracy: 0.4966
25 Epoch 13/15
26 782/782 [=====] - 4s 5ms/step - loss: 1.3774 - accuracy: 0.5057 - val_loss:
   1.3549 - val_accuracy: 0.5232
27 Epoch 14/15
28 782/782 [=====] - 4s 5ms/step - loss: 1.3663 - accuracy: 0.5114 - val_loss:
   1.3897 - val_accuracy: 0.5019
29 Epoch 15/15
30 782/782 [=====] - 4s 5ms/step - loss: 1.3532 - accuracy: 0.5146 - val_loss:
   1.3468 - val_accuracy: 0.5167
31
32

```

Figure 46: Adam Optimizer

2.5.2 RMSprop

```

1 Epoch 1/15
2 782/782 [=====] - 4s 5ms/step - loss: 1.9534 - accuracy: 0.2977 - val_loss:
   1.8734 - val_accuracy: 0.3299
3 Epoch 2/15
4 782/782 [=====] - 4s 5ms/step - loss: 1.7362 - accuracy: 0.3749 - val_loss:
   1.8528 - val_accuracy: 0.3441
5 Epoch 3/15
6 782/782 [=====] - 4s 5ms/step - loss: 1.6628 - accuracy: 0.4033 - val_loss:
   1.6178 - val_accuracy: 0.4163
7 Epoch 4/15
8 782/782 [=====] - 4s 5ms/step - loss: 1.6131 - accuracy: 0.4245 - val_loss:
   1.7880 - val_accuracy: 0.3693
9 Epoch 5/15
10 782/782 [=====] - 4s 5ms/step - loss: 1.5732 - accuracy: 0.4409 - val_loss:
   1.5200 - val_accuracy: 0.4563
11 Epoch 6/15
12 782/782 [=====] - 4s 5ms/step - loss: 1.5441 - accuracy: 0.4479 - val_loss:
   1.6083 - val_accuracy: 0.4284
13 Epoch 7/15
14 782/782 [=====] - 3s 4ms/step - loss: 1.5176 - accuracy: 0.4593 - val_loss:
   1.5595 - val_accuracy: 0.4421
15 Epoch 8/15
16 782/782 [=====] - 4s 5ms/step - loss: 1.4970 - accuracy: 0.4672 - val_loss:
   1.5623 - val_accuracy: 0.4330
17 Epoch 9/15
18 782/782 [=====] - 4s 5ms/step - loss: 1.4840 - accuracy: 0.4732 - val_loss:
   1.5004 - val_accuracy: 0.4614
19 Epoch 10/15
20 782/782 [=====] - 3s 4ms/step - loss: 1.4679 - accuracy: 0.4770 - val_loss:
   1.4534 - val_accuracy: 0.4788
21 Epoch 11/15
22 782/782 [=====] - 4s 5ms/step - loss: 1.4498 - accuracy: 0.4851 - val_loss:
   1.5108 - val_accuracy: 0.4575
23 Epoch 12/15
24 782/782 [=====] - 4s 5ms/step - loss: 1.4394 - accuracy: 0.4883 - val_loss:
   1.4865 - val_accuracy: 0.4785
25 Epoch 13/15
26 782/782 [=====] - 4s 5ms/step - loss: 1.4268 - accuracy: 0.4922 - val_loss:
   1.6755 - val_accuracy: 0.4242
27 Epoch 14/15
28 782/782 [=====] - 4s 5ms/step - loss: 1.4184 - accuracy: 0.4972 - val_loss:
   1.3920 - val_accuracy: 0.5095
29 Epoch 15/15
30 782/782 [=====] - 3s 4ms/step - loss: 1.4104 - accuracy: 0.4990 - val_loss:
   1.5326 - val_accuracy: 0.4620
31
32

```

Figure 47: RMSprop Optimizer

As we can see the **Adam** optimizer is a little better than the other ones, this is an SGD based optimizer which is frequently used in neural networks.

2.6 Part 5 : The best model

After studying the previous parts, we conclude that using **Relu** and **Softmax** activations, the **categorical_crossentropy** loss function and the **Adam** optimizer yields to the best networks.

We shall go on to train this model and get the accuracy for the test data. After this we shall check its precision, recall and f-score.

The code to train is as follows.

```
1 X_train, Y_train, X_test, Y_test = Loader()
2 X_train, Y_train, X_test, Y_test, X_valid, Y_valid =
  preprocessing(X_train, Y_train, X_test, Y_test)
3 model64adam = NeuralNetwork('relu', 'relu', 'softmax', '
  categorical_crossentropy', 'adam')
4 history64adam = model64adam.fit(X_train, Y_train, epochs=15,
  batch_size=64, validation_data=(X_valid, Y_valid), verbose=1)
5
```

Figure 48: Code to train

```
1 1563/1563 [=====] - 2s 1ms/step
2 1563/1563 [=====] - 3s 2ms/step -
  loss: 1.3097 - accuracy: 0.5346
3 test loss = 1.309681
4 test accuracy = 0.534580
5 Accuracy: 0.534580
6 Precision: 0.535605
7 Recall: 0.534580
8 F1 score: 0.531291
9
```

Figure 49: Complete Analysis

```

1 Epoch 1/15
2 782/782 [=====] - 5s 5ms/step - loss: 1.8566 - accuracy: 0.3291 - val_loss:
   1.7294 - val_accuracy: 0.3751
3 Epoch 2/15
4 782/782 [=====] - 4s 5ms/step - loss: 1.6843 - accuracy: 0.3971 - val_loss:
   1.6944 - val_accuracy: 0.3895
5 Epoch 3/15
6 782/782 [=====] - 4s 5ms/step - loss: 1.6147 - accuracy: 0.4220 - val_loss:
   1.6166 - val_accuracy: 0.4195
7 Epoch 4/15
8 782/782 [=====] - 4s 5ms/step - loss: 1.5671 - accuracy: 0.4405 - val_loss:
   1.5939 - val_accuracy: 0.4327
9 Epoch 5/15
10 782/782 [=====] - 4s 5ms/step - loss: 1.5293 - accuracy: 0.4526 - val_loss:
   1.5435 - val_accuracy: 0.4489
11 Epoch 6/15
12 782/782 [=====] - 4s 5ms/step - loss: 1.4946 - accuracy: 0.4676 - val_loss:
   1.4566 - val_accuracy: 0.4780
13 Epoch 7/15
14 782/782 [=====] - 4s 5ms/step - loss: 1.4755 - accuracy: 0.4745 - val_loss:
   1.4519 - val_accuracy: 0.4825
15 Epoch 8/15
16 782/782 [=====] - 4s 5ms/step - loss: 1.4519 - accuracy: 0.4811 - val_loss:
   1.4569 - val_accuracy: 0.4805
17 Epoch 9/15
18 782/782 [=====] - 4s 5ms/step - loss: 1.4325 - accuracy: 0.4913 - val_loss:
   1.4046 - val_accuracy: 0.4980
19 Epoch 10/15
20 782/782 [=====] - 4s 6ms/step - loss: 1.4058 - accuracy: 0.4974 - val_loss:
   1.3695 - val_accuracy: 0.5075
21 Epoch 11/15
22 782/782 [=====] - 4s 5ms/step - loss: 1.3995 - accuracy: 0.5028 - val_loss:
   1.3467 - val_accuracy: 0.5180
23 Epoch 12/15
24 782/782 [=====] - 4s 5ms/step - loss: 1.3784 - accuracy: 0.5091 - val_loss:
   1.3615 - val_accuracy: 0.5125
25 Epoch 13/15
26 782/782 [=====] - 4s 5ms/step - loss: 1.3757 - accuracy: 0.5114 - val_loss:
   1.3614 - val_accuracy: 0.5222
27 Epoch 14/15
28 782/782 [=====] - 4s 5ms/step - loss: 1.3526 - accuracy: 0.5187 - val_loss:
   1.3659 - val_accuracy: 0.5170
29 Epoch 15/15
30 782/782 [=====] - 4s 5ms/step - loss: 1.3420 - accuracy: 0.5242 - val_loss:
   1.3011 - val_accuracy: 0.5394
31
32

```

Figure 50: The best model

Now we shall take a look at the confusion matrix.

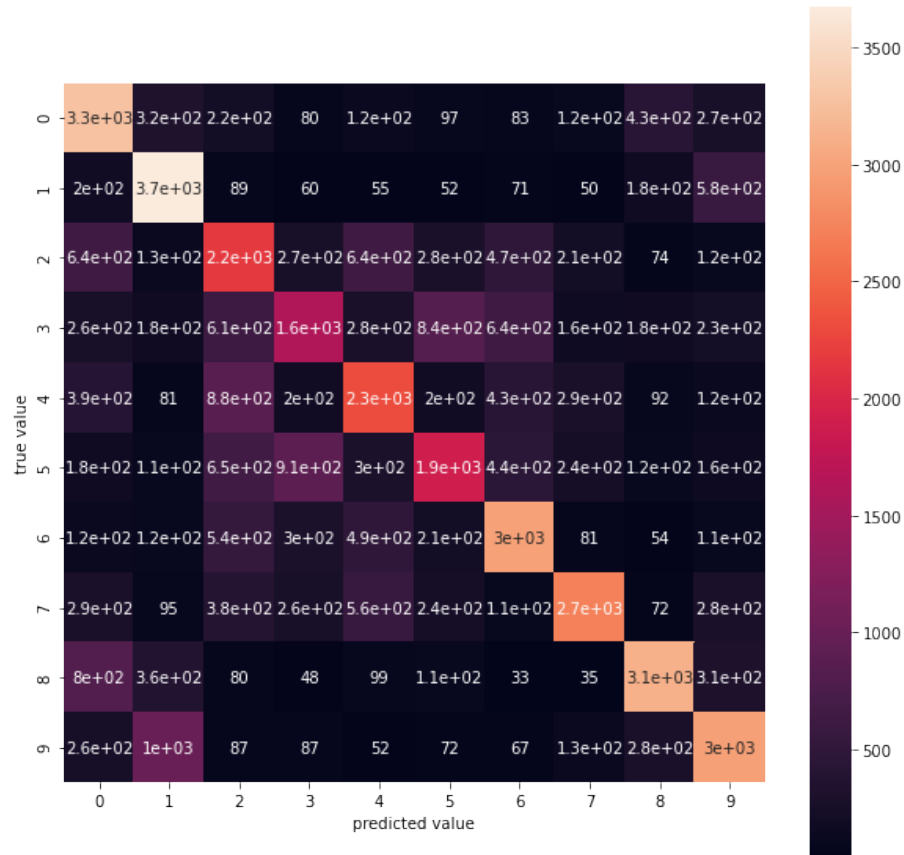


Figure 51: Confusion Matrix

This concludes the analysis of the **MLP**, next we shall move on to the **MLP+CNN** model.

MLP+CNN

First of all we shall include our new preprocessing method.

```

1  def preprocessingcnn(X_train, Y_train, X_test, Y_test):
2      # Transform label indices to one-hot encoded vectors
3      Y_train = to_categorical(Y_train, num_classes = len(Cifar_Labels))
4      Y_test = to_categorical(Y_test, num_classes = len(Cifar_Labels))
5      #Because working with 4 dimensions is not possible but is possible in cnns, we
6      #change the dimensions
7      # X_train = np.reshape(X_train, (X_train.shape[0], 32*32*3))
8      #X_test = np.reshape(X_test, (X_test.shape[0], 32*32*3))
9      X_train = X_train.astype('float32')
10     X_test = X_test.astype('float32')
11     #Here we perform normalization
12     X_train = X_train/255
13     X_test = X_test/255
14     #Test data must be kept unseen from the model until final evaluation hence we need
15     #validation dataset.
16     _, X_valid, _, Y_valid = train_test_split(X_train, Y_train,
17     test_size=0.2, random_state= 8)
18     return X_train, Y_train, X_test, Y_test, X_valid, Y_valid
19

```

Figure 52: Preprocessing

2.7 Part 1 (CNN) : The preliminary CNN

The designed network is depicted below.

```

1  def NeuralNetworkcnn(1, opt):
2      model = Sequential()
3      model.add(Conv2D(32, (3, 3), padding='same', input_shape=(32,
4      32, 3), activation="relu"))
5      model.add(Conv2D(32, (3, 3), activation="relu"))
6      model.add(Flatten())
7      model.add(Dense(128, activation='relu', input_dim = 32*32*3))
8      model.add(Dense(128, activation='relu'))
9      model.add(Dense(10, activation="softmax"))
10     model.compile(loss=1, optimizer=opt, metrics=['accuracy'])
11     return model

```

Figure 53: CNN

The results are better than the **MLP** with about 10% more test accuracy and less error.

2.8 Part 2 (CNN) : Pooling and Batch Normalization

Firstly I shall explain the batch normalization and max pooling layers in convolutional neural networks.

2.8.1 Batch Normalization

Batch normalization is a normalization technique, this technique is utilized between the layers of a neural network instead of normalizing the raw data, batch normalization is done along mini-batches instead of the entire dataset. It is useful for speeding up training and using higher learning rates, making learning easier. It can be formulated as below.

$$z^N = \left(\frac{z - m_z}{s_z} \right), \quad s_z : \text{Standard Deviation}$$

2.8.2 Pooling

We know that the pooling operation involves sliding a two-dimensional filter over each channel of feature map and summarising the features lying within the region covered by the filter.

Pooling layers are used to reduce the dimensions of feature maps. So pooling reduces the number of parameters to learn and the amount of computation needed.

Also it is important to imply that pooling summarises the features present in a region of the feature map which is created by the CNN.

There many types of pooling such as max pooling, average pooling, global pooling etc...

In this project I have used max pooling.

The designed network is depicted below.

```
1 def NeuralNetworkcnnpoolbatchnorm(1,opt):
2     model = Sequential()
3     model.add(Conv2D(32, (3, 3), padding='same', input_shape=(32,
4         32, 3), activation="relu"))
5     model.add(BatchNormalization())
6     model.add(MaxPooling2D(pool_size=(2, 2)))
7     model.add(Conv2D(32, (3, 3), activation="relu"))
8     model.add(BatchNormalization())
9     model.add(MaxPooling2D(pool_size=(2, 2)))
10    model.add(Flatten())
11    model.add(BatchNormalization())
12    model.add(Dense(128,activation='relu',input_dim = 32*32*3))
13    model.add(Dense(128,activation='relu'))
14    model.add(Dense(10, activation="softmax"))
15    model.compile(loss=l, optimizer=opt, metrics=['accuracy'])
16    return model
```

Figure 54: CNN with max pooling and batch normalization

The code to train and results are as follows.

```
1 X_train, Y_train, X_test, Y_test = Loader()
2 X_train, Y_train, X_test, Y_test, X_valid, Y_valid =
3     preprocessingcnn(X_train, Y_train, X_test, Y_test)
4     model64adamcnnbp = NeuralNetworkcnnpoolbatchnorm('
5         categorical_crossentropy', 'adam')
6     history64adamcnnbp = model64adamcnnbp.fit(X_train, Y_train,
7         epochs=15, batch_size=64, validation_data=(X_valid, Y_valid),
8         verbose=1)
```

Figure 55: Code to train

Train Results

```

1 Epoch 1/15
2 782/782 [=====] - 31s 38ms/step - loss: 1.2756 - accuracy: 0.5460 - val_loss:
3   0.9410 - val_accuracy: 0.6702
4 Epoch 2/15
5 782/782 [=====] - 30s 39ms/step - loss: 0.9146 - accuracy: 0.6784 - val_loss:
6   0.8208 - val_accuracy: 0.7046
7 Epoch 3/15
8 782/782 [=====] - 29s 37ms/step - loss: 0.7608 - accuracy: 0.7303 - val_loss:
9   0.6284 - val_accuracy: 0.7791
10 Epoch 4/15
11 782/782 [=====] - 30s 38ms/step - loss: 0.6422 - accuracy: 0.7708 - val_loss:
12   0.5646 - val_accuracy: 0.7972
13 Epoch 5/15
14 782/782 [=====] - 29s 37ms/step - loss: 0.5513 - accuracy: 0.8042 - val_loss:
15   0.4384 - val_accuracy: 0.8470
16 Epoch 6/15
17 782/782 [=====] - 29s 37ms/step - loss: 0.4718 - accuracy: 0.8312 - val_loss:
18   0.3684 - val_accuracy: 0.8715
19 Epoch 7/15
20 782/782 [=====] - 30s 38ms/step - loss: 0.3989 - accuracy: 0.8580 - val_loss:
21   0.4330 - val_accuracy: 0.8453
22 Epoch 8/15
23 782/782 [=====] - 31s 40ms/step - loss: 0.3500 - accuracy: 0.8758 - val_loss:
24   0.2699 - val_accuracy: 0.9062
25 Epoch 9/15
26 782/782 [=====] - 30s 38ms/step - loss: 0.3027 - accuracy: 0.8913 - val_loss:
27   0.3269 - val_accuracy: 0.8838
28 Epoch 10/15
29 782/782 [=====] - 29s 37ms/step - loss: 0.2583 - accuracy: 0.9075 - val_loss:
30   0.2057 - val_accuracy: 0.9303
31 Epoch 11/15
32 782/782 [=====] - 29s 37ms/step - loss: 0.2347 - accuracy: 0.9154 - val_loss:
33   0.3131 - val_accuracy: 0.8915
34 Epoch 12/15
35 782/782 [=====] - 29s 37ms/step - loss: 0.2066 - accuracy: 0.9281 - val_loss:
36   0.1484 - val_accuracy: 0.9482
37 Epoch 13/15
38 782/782 [=====] - 30s 38ms/step - loss: 0.1833 - accuracy: 0.9344 - val_loss:
39   0.1953 - val_accuracy: 0.9312
40 Epoch 14/15
41 782/782 [=====] - 30s 38ms/step - loss: 0.1777 - accuracy: 0.9371 - val_loss:
42   0.1594 - val_accuracy: 0.9443
43 Epoch 15/15
44 782/782 [=====] - 30s 38ms/step - loss: 0.1610 - accuracy: 0.9422 - val_loss:
45   0.1982 - val_accuracy: 0.9314

```

Figure 56: CNN with max pooling and batch normalization

Accuracy Plot




Figure 57: Accuracy Plot

Loss Plot




Figure 58: Loss Plot

Confusion matrix and accuracy and error

```

1  313/313 [=====] - 2s 6ms/step
2  313/313 [=====] - 2s 6ms/step - loss :
   1.6031 - accuracy: 0.6761
3  test loss = 1.603063
4  test accuracy = 0.676100
5  Accuracy: 0.676100
6  Precision: 0.700380
7  Recall: 0.676100
8  F1 score: 0.681632
9

```

Figure 59: Accuracy and error for CNN with max pooling and batch normalization




Figure 60: Confusion Matrix

As we can see both the loss and accuracy has improved compared to the MLP architecture.

2.9 Part 3 (CNN) : Adding dropout

Firstly I shall give some explanation about dropout and what it is.

2.9.1 Dropout

As instructed in the lectures, the term dropout means dropping out the nodes in a neural network, which can be the input or hidden layer nodes. All the forward and backward connections are removed temporarily so we generally create a new architecture.

Dropout is very useful in solving the overfitting problem, we know that in overfitting neurons might change in a way that fixes the mistakes of other neurons, dropout prevents neurons fixing the mistakes of other neurons, hence by randomly dropping a few nodes, it forces the layers to take more or less responsibility for the input by taking a probabilistic approach.

So to conclude our explanation we say that we use dropout to fix the overfitting problem.

The designed network is depicted below.

```

1  def NeuralNetworkcnnndrop(1,opt):
2  model = Sequential()
3  model.add(Conv2D(32, (3, 3), padding='same', input_shape=(32,
    32, 3), activation="relu"))
4  model.add(BatchNormalization())
5  model.add(MaxPooling2D(pool_size=(2, 2)))
6  model.add(Dropout(0.25))
7  model.add(Conv2D(32, (3, 3), activation="relu"))
8  model.add(BatchNormalization())
9  model.add(MaxPooling2D(pool_size=(2, 2)))
10 model.add(Dropout(0.25))
11 model.add(Flatten())
12 model.add(BatchNormalization())
13 model.add(Dropout(0.25))
14 model.add(Dense(128,activation='relu',input_dim = 32*32*3))
15 model.add(Dense(128,activation='relu'))
16 model.add(Dense(10, activation="softmax"))
17 model.compile(loss=1, optimizer=opt, metrics=['accuracy'])
18 return model
19

```

Figure 61: CNN with max pooling and batch normalization and dropout

The code to train and results are as follows.

```

1  X_train, Y_train, X_test, Y_test = Loader()
2  X_train, Y_train, X_test, Y_test, X_valid, Y_valid = preprocessingcnn(X_train, Y_train, X_test, Y_test)
3  model64adamcnnbpd = NeuralNetworkcnnndrop('categorical_crossentropy', 'adam')
4  history64adamcnnbpd = model64adamcnnbpd.fit(X_train, Y_train, epochs=15, batch_size=64, validation_data=(
    X_valid, Y_valid), verbose=1)
5

```

Figure 62: Code to train

Train Results

```

1 Epoch 1/15
2 782/782 [=====] - 37s 45ms/step - loss: 1.4404 - accuracy: 0.4823 - val_loss:
   1.3081 - val_accuracy: 0.5348
3 Epoch 2/15
4 782/782 [=====] - 34s 44ms/step - loss: 1.1163 - accuracy: 0.6028 - val_loss:
   1.2351 - val_accuracy: 0.5748
5 Epoch 3/15
6 782/782 [=====] - 34s 44ms/step - loss: 0.9943 - accuracy: 0.6451 - val_loss:
   0.7949 - val_accuracy: 0.7186
7 Epoch 4/15
8 782/782 [=====] - 34s 43ms/step - loss: 0.9184 - accuracy: 0.6743 - val_loss:
   0.8778 - val_accuracy: 0.6910
9 Epoch 5/15
10 782/782 [=====] - 35s 45ms/step - loss: 0.8606 - accuracy: 0.6938 - val_loss:
    0.7139 - val_accuracy: 0.7553
11 Epoch 6/15
12 782/782 [=====] - 33s 42ms/step - loss: 0.8170 - accuracy: 0.7138 - val_loss:
    0.6629 - val_accuracy: 0.7566
13 Epoch 7/15
14 782/782 [=====] - 34s 43ms/step - loss: 0.7748 - accuracy: 0.7252 - val_loss:
    0.6814 - val_accuracy: 0.7533
15 Epoch 8/15
16 782/782 [=====] - 33s 42ms/step - loss: 0.7454 - accuracy: 0.7350 - val_loss:
    0.7308 - val_accuracy: 0.7425
17 Epoch 9/15
18 782/782 [=====] - 33s 42ms/step - loss: 0.7265 - accuracy: 0.7419 - val_loss:
    0.6277 - val_accuracy: 0.7802
19 Epoch 10/15
20 782/782 [=====] - 33s 42ms/step - loss: 0.7031 - accuracy: 0.7512 - val_loss:
    0.5166 - val_accuracy: 0.8168
21 Epoch 11/15
22 782/782 [=====] - 33s 42ms/step - loss: 0.6795 - accuracy: 0.7598 - val_loss:
    0.4891 - val_accuracy: 0.8314
23 Epoch 12/15
24 782/782 [=====] - 34s 44ms/step - loss: 0.6637 - accuracy: 0.7643 - val_loss:
    0.4912 - val_accuracy: 0.8278
25 Epoch 13/15
26 782/782 [=====] - 34s 43ms/step - loss: 0.6470 - accuracy: 0.7701 - val_loss:
    0.4605 - val_accuracy: 0.8353
27 Epoch 14/15
28 782/782 [=====] - 34s 44ms/step - loss: 0.6355 - accuracy: 0.7740 - val_loss:
    0.4919 - val_accuracy: 0.8290
29 Epoch 15/15
30 782/782 [=====] - 34s 44ms/step - loss: 0.6223 - accuracy: 0.7798 - val_loss:
    0.4260 - val_accuracy: 0.8478
31

```

Figure 63: CNN with max pooling and batch normalization and dropout

Accuracy Plot

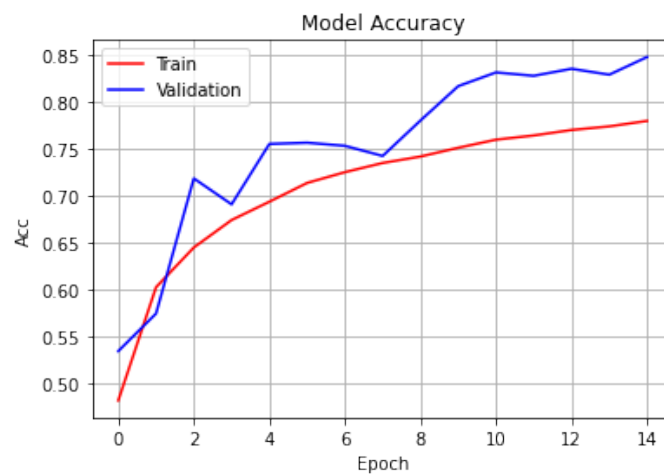


Figure 64: Accuracy Plot

Loss Plot




Figure 65: Loss Plot

Confusion matrix and accuracy and error

```

1  313/313 [=====] - 2s 6ms/step
2  313/313 [=====] - 2s 6ms/step - loss:
    0.7780 - accuracy: 0.7298
3  test loss = 0.778038
4  test accuracy = 0.729800
5  Accuracy: 0.729800
6  Precision: 0.747035
7  Recall: 0.729800
8  F1 score: 0.731516
9

```

Figure 66: Accuracy and error for CNN with max pooling and batch normalization




Figure 67: Confusion Matrix

We see that the overfitting problem is fixed and we get 72.98% accuracy.

2.10 Part 4 (CNN) : Early Stopping

Here I shall first explain early stopping.

2.10.1 Early Stopping

As depicted in every neural networks course, early stopping is a technique used for regularization for neural networks that stops training on a certain epoch when the parameter updates no longer begin to yield improves on our validation set.

Metrics such is minimum changes in the monitored quantity and number of epochs with no improvement and baseline values for the monitored quantity can all be metrics in early stopping a neural network.

A very important graph in early stopping is shown below.




Figure 68: Early stopping

Now we shall implement early stopping as the final part of this section of the assignment.

The designed network is depicted below.

```

1  def NeuralNetworkcnnEarlyStop(1,opt):
2  model = Sequential()
3  model.add(Conv2D(32, (3, 3), padding='same', input_shape=(32,
    32, 3), activation="relu"))
4  model.add(BatchNormalization())
5  model.add(MaxPooling2D(pool_size=(2, 2)))
6  model.add(Dropout(0.25))
7  model.add(Conv2D(32, (3, 3), activation="relu"))
8  model.add(BatchNormalization())
9  model.add(MaxPooling2D(pool_size=(2, 2)))
10 model.add(Dropout(0.25))
11 model.add(Flatten())
12 model.add(BatchNormalization())
13 model.add(Dropout(0.25))
14 model.add(Dense(128,activation='relu',input_dim = 32*32*3))
15 model.add(Dense(128,activation='relu'))
16 model.add(Dense(10, activation="softmax"))
17 model.compile(loss=1, optimizer=opt, metrics=['accuracy'])
18 return model
19

```

Figure 69: CNN Early Stopping

The code to train and results are as follows.

```

1  X_train, Y_train, X_test, Y_test = Loader()
2  X_train, Y_train, X_test, Y_test, X_valid, Y_valid = preprocessingcnn(X_train, Y_train, X_test, Y_test)
3  model64adamcnnr1 = NeuralNetworkcnnEarlyStop('categorical_crossentropy', 'adam')
4  es = EarlyStopping(monitor='val_accuracy', mode='max', min_delta=1, verbose=1, patience=5)
5  history64adamcnnr1 = model64adamcnnr1.fit(X_train, Y_train, epochs=15, batch_size=64, validation_data=(
    X_valid, Y_valid), callbacks=[es], verbose=1)
6

```

Figure 70: Code to train

Train Results

```

1 Epoch 1/15
2 782/782 [=====] - 35s 43ms/step - loss: 1.4765 - accuracy: 0.4707 - val_loss:
   1.3721 - val_accuracy: 0.5186
3 Epoch 2/15
4 782/782 [=====] - 33s 43ms/step - loss: 1.1906 - accuracy: 0.5742 - val_loss:
   1.3616 - val_accuracy: 0.5255
5 Epoch 3/15
6 782/782 [=====] - 36s 45ms/step - loss: 1.0452 - accuracy: 0.6292 - val_loss:
   1.1625 - val_accuracy: 0.6078
7 Epoch 4/15
8 782/782 [=====] - 37s 47ms/step - loss: 0.9614 - accuracy: 0.6553 - val_loss:
   0.8053 - val_accuracy: 0.7223
9 Epoch 5/15
10 782/782 [=====] - 34s 43ms/step - loss: 0.8970 - accuracy: 0.6813 - val_loss:
   0.8304 - val_accuracy: 0.7072
11 Epoch 6/15
12 782/782 [=====] - 34s 43ms/step - loss: 0.8488 - accuracy: 0.6985 - val_loss:
   0.6863 - val_accuracy: 0.7601
13 Epoch 6: early stopping
14

```

Figure 71: CNN Early Stopping

Accuracy Plot

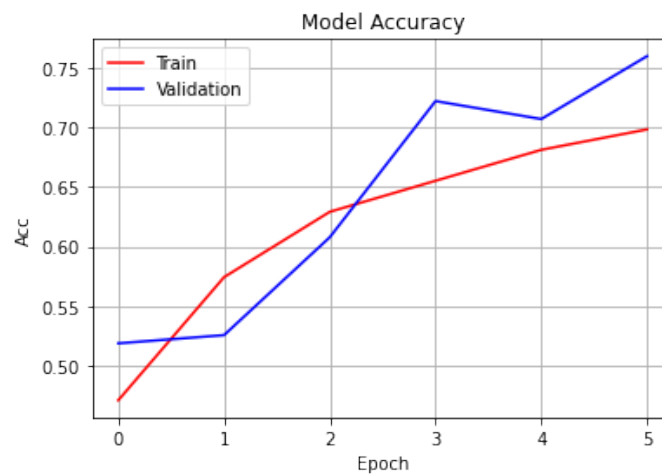


Figure 72: Accuracy Plot

Loss Plot




Figure 73: Loss Plot

Confusion matrix and accuracy and error

```

1  313/313 [=====] - 2s 6ms/step
2  313/313 [=====] - 2s 6ms/step - loss:
   0.8985 - accuracy: 0.6888
3  test loss = 0.898465
4  test accuracy = 0.688800
5  Accuracy: 0.688800
6  Precision: 0.711041
7  Recall: 0.688800
8  F1 score: 0.692636
9

```

Figure 74: Accuracy and error for CNN early stopping




Figure 75: Confusion Matrix

3 Question 3 : Transfer Learning for EfficientNet network

3.1 Part A : Studying EfficientNet

Here we shall give some preliminary explanation about convolutional neural networks to build up our explanation for the EfficientNet architecture.

We know CNNs are normally developed at a fixed resource budget, and then scaled up to achieve better accuracy if more resources are available. In the EfficientNet paper a new scaling method that uniformly scales all dimensions of depth/width/resolution using a simple yet highly effective compound coefficient is proposed.

It is of common knowledge that scaling up CNNs is widely used to get better accuracy for example, ResNet can be scaled up from ResNet-18 to ResNet-200 by using more layers.

We have many ways to scale up CNNs, the most common way is to perform scaling by depth. In the given paper the main question is that is there a principled method to scale up ConvNets that can achieve better accuracy and efficiency? We can see some scaling methods in the following picture.

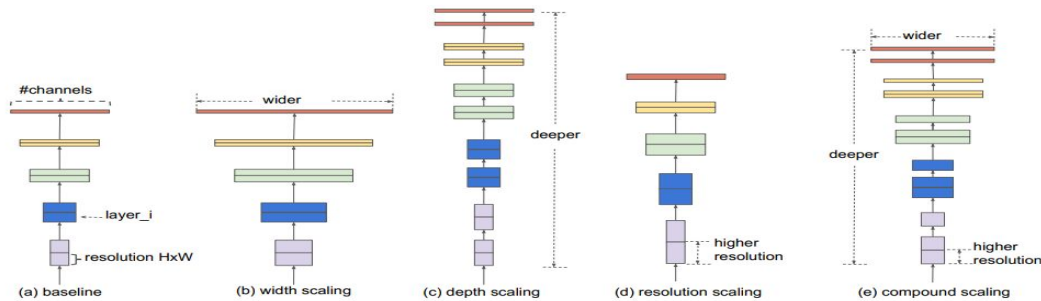


Figure 76: Different scaling models

3.1.1 Network Architecture

As explained in the paper due to the fact that model scaling does not change layer operators in baseline network, having a good baseline network is crucial, a new mobile-size baseline, called EfficientNet has been developed.

The baseline network has been developed by leveraging a multi-objective neural architecture search that optimizes both accuracy and FLOPS(Floating Point Operations Per Second). We optimize FLOPS rather than latency since we are not targeting any specific hardware device. This search produces an efficient network, which is named EfficientNet-B0.

We can see the EfficientNet-B0 baseline network table and structure below.

Stage i	Operator \mathcal{F}_i	Resolution $\hat{H}_i \times \hat{W}_i$	#Channels \hat{C}_i	#Layers \hat{L}_i
1	Conv3x3	224×224	32	1
2	MBConv1, k3x3	112×112	16	1
3	MBConv6, k3x3	112×112	24	2
4	MBConv6, k5x5	56×56	40	2
5	MBConv6, k3x3	28×28	80	3
6	MBConv6, k5x5	28×28	112	3
7	MBConv6, k5x5	14×14	192	4
8	MBConv6, k3x3	7×7	320	1
9	Conv1x1 & Pooling & FC	7×7	1280	1

Figure 77: EfficientNet-B0 baseline network




Figure 78: EfficientNet-B0 baseline network

3.1.2 Different EfficientNet Architectures

We must pay sincere attention that the **EfficientNetB0** model is the baseline, a full family of EfficientNet model have been developed based on **EfficientNetB0**, these models are named from **EfficientNetB1** to **EfficientNetB7**.

In the following photos we can see the difference between the architectures of these models.

EfficientNetB0

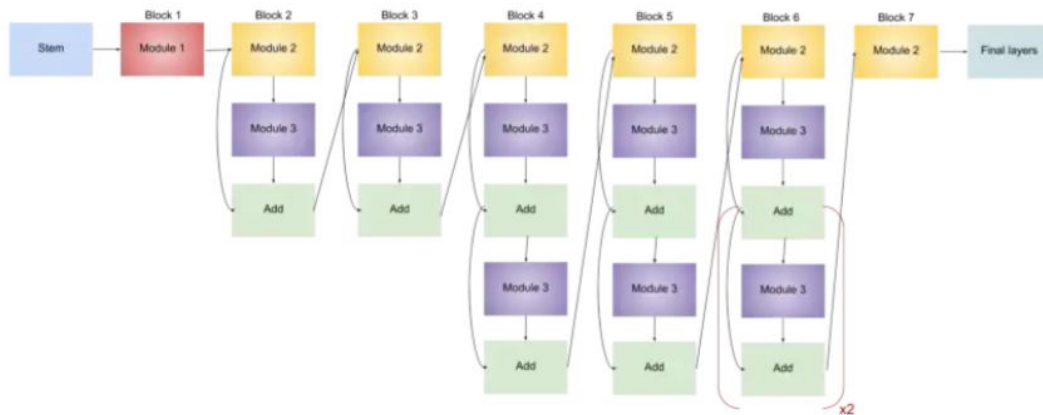


Figure 79: EfficientNet-B0 model

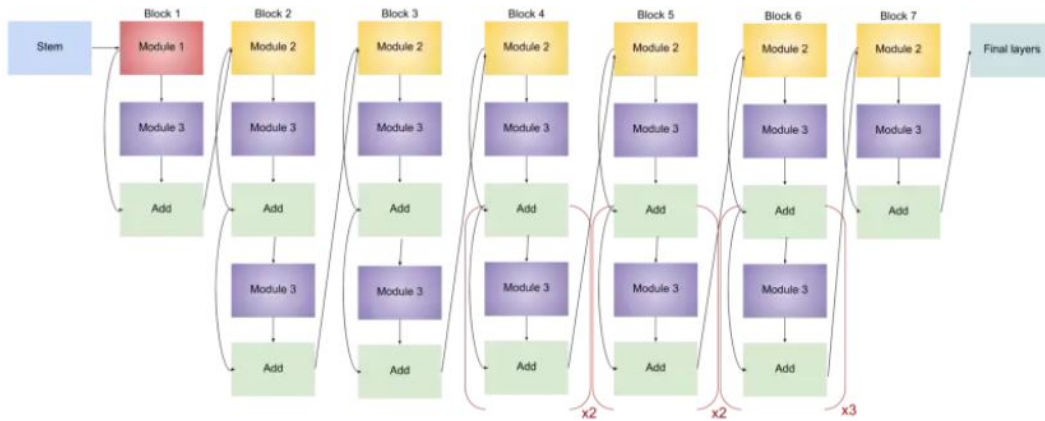
EfficientNetB1

Figure 80: EfficientNet-B1 model

EfficientNetB2

Exactly the same as B1 but with more parameters.

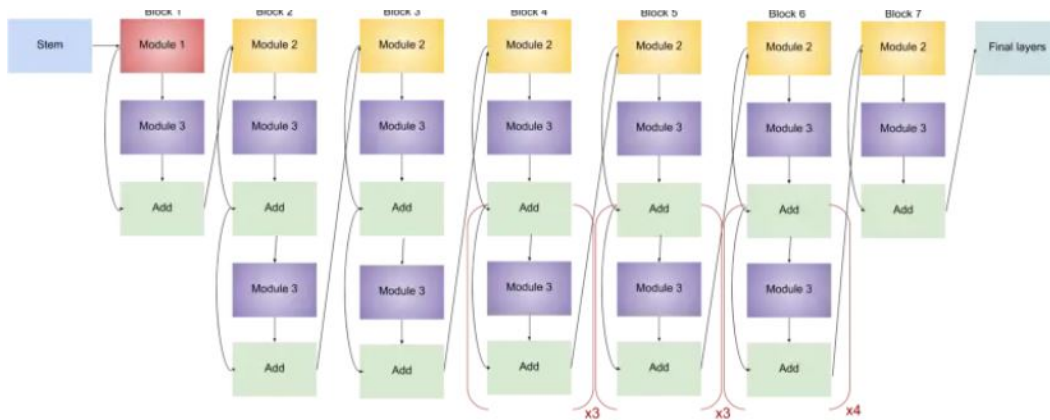
EfficientNetB3

Figure 81: EfficientNet-B3 model

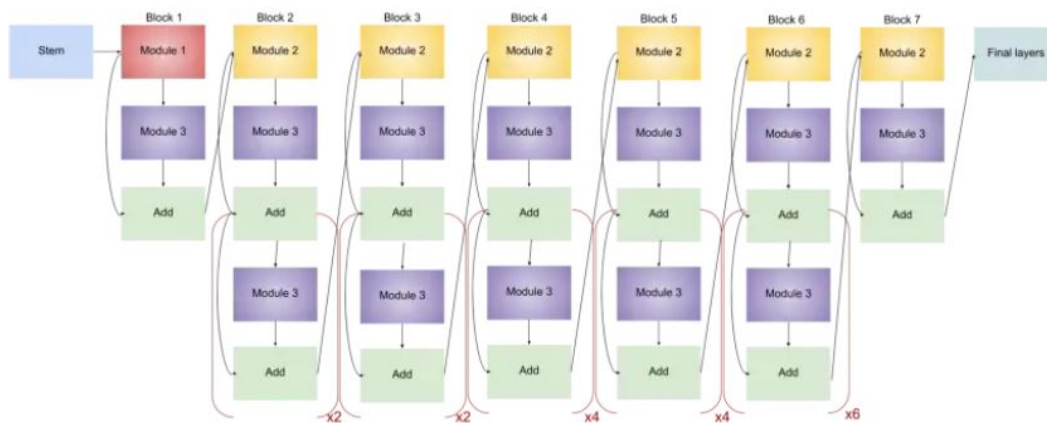
EfficientNetB4

Figure 82: EfficientNet-B4 model

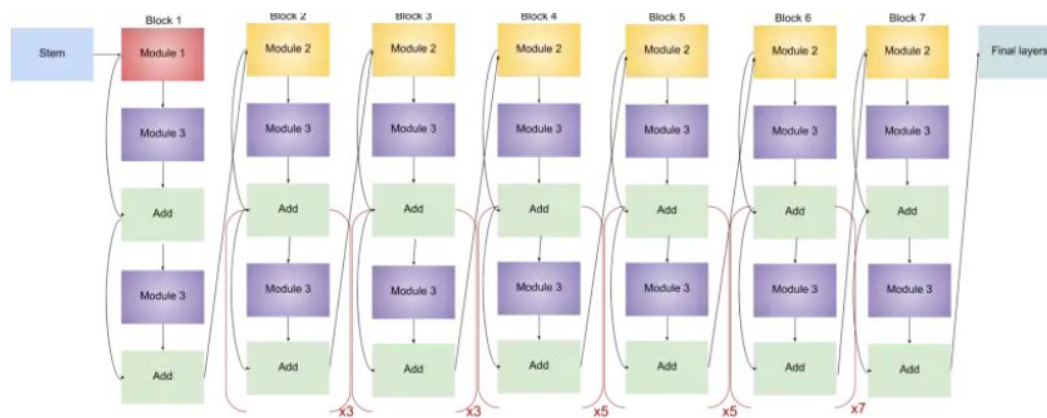
EfficientNetB5

Figure 83: EfficientNet-B5 model

EfficientNetB6

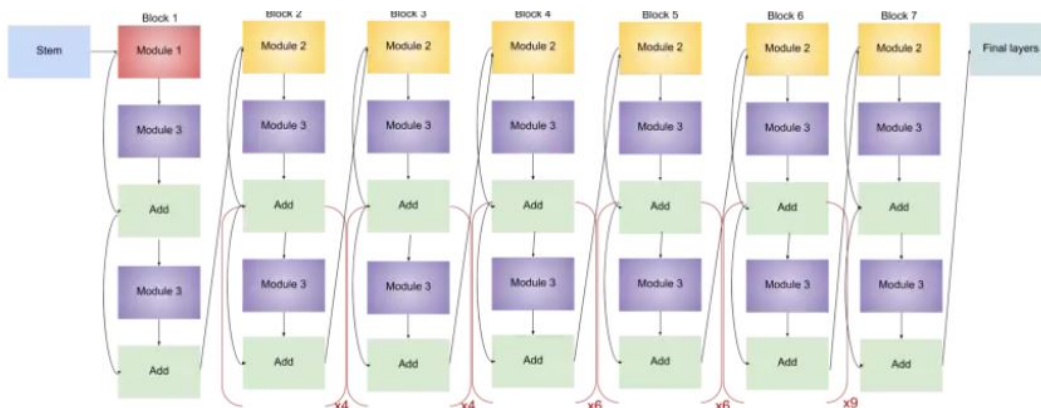


Figure 84: EfficientNet-B6 model

EfficientNetB7

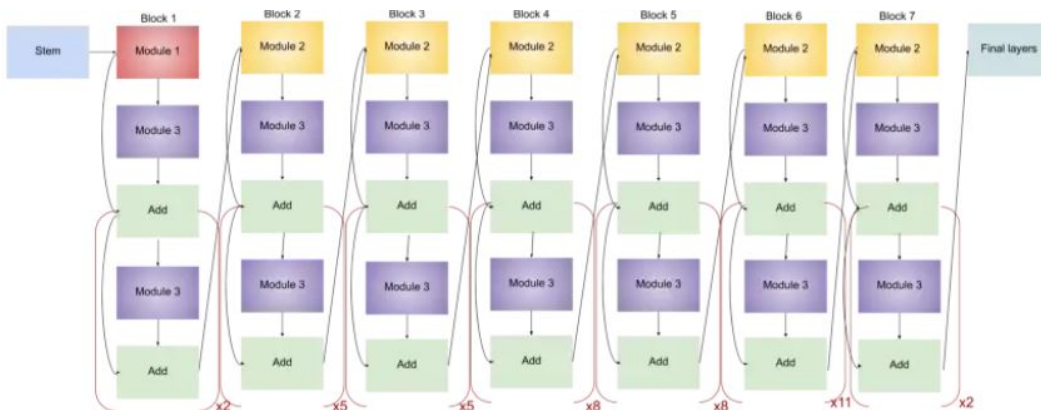


Figure 85: EfficientNet-B7 model

The main difference is in the number of channels which is depicted in the table below.

Stage	B1	B2	B3	B4	B5	B6	B7
1	32	32	40	48	48	56	64
2	16	16	24	24	24	32	32
3	24	24	32	32	40	40	48
4	40	48	48	56	64	72	80
5	80	88	96	112	128	144	160
6	112	120	136	160	176	200	224
7	192	208	232	272	304	344	384
8	320	352	384	448	512	576	640
9	1280	1408	1536	1792	2048	2304	2560

Figure 86: EfficientNet models

In the next table we can see the EfficientNet performance results on ImageNet.

Model	Top-1 Acc.	Top-5 Acc.	#Params	Ratio-to-EfficientNet	#FLOPS	Ratio-to-EfficientNet
EfficientNet-B0	76.3%	93.2%	5.3M	1x	0.39B	1x
ResNet-50 (He et al., 2016)	76.0%	93.0%	26M	4.9x	4.1B	11x
DenseNet-169 (Huang et al., 2017)	76.2%	93.2%	14M	2.6x	3.5B	8.9x
EfficientNet-B1	78.8%	94.4%	7.8M	1x	0.70B	1x
ResNet-152 (He et al., 2016)	77.8%	93.8%	60M	7.6x	11B	16x
DenseNet-264 (Huang et al., 2017)	77.9%	93.9%	34M	4.3x	6.0B	8.6x
Inception-v3 (Szegedy et al., 2016)	78.8%	94.4%	24M	3.0x	5.7B	8.1x
Xception (Chollet, 2017)	79.0%	94.5%	23M	3.0x	8.4B	12x
EfficientNet-B2	79.8%	94.9%	9.2M	1x	1.0B	1x
Inception-v4 (Szegedy et al., 2017)	80.0%	95.0%	48M	5.2x	13B	13x
Inception-resnet-v2 (Szegedy et al., 2017)	80.1%	95.1%	56M	6.1x	13B	13x
EfficientNet-B3	81.1%	95.5%	12M	1x	1.8B	1x
ResNeXt-101 (Xie et al., 2017)	80.9%	95.6%	84M	7.0x	32B	18x
PolyNet (Zhang et al., 2017)	81.3%	95.8%	92M	7.7x	35B	19x
EfficientNet-B4	82.6%	96.3%	19M	1x	4.2B	1x
SENet (Hu et al., 2018)	82.7%	96.2%	146M	7.7x	42B	10x
NASNet-A (Zoph et al., 2018)	82.7%	96.2%	89M	4.7x	24B	5.7x
AmoebaNet-A (Real et al., 2019)	82.8%	96.1%	87M	4.6x	23B	5.5x
PNASNet (Liu et al., 2018)	82.9%	96.2%	86M	4.5x	23B	6.0x
EfficientNet-B5	83.3%	96.7%	30M	1x	9.9B	1x
AmoebaNet-C (Cubuk et al., 2019)	83.5%	96.5%	155M	5.2x	41B	4.1x
EfficientNet-B6	84.0%	96.9%	43M	1x	19B	1x
EfficientNet-B7	84.4%	97.1%	66M	1x	37B	1x
GPipe (Huang et al., 2018)	84.3%	97.0%	557M	8.4x	-	-

Figure 87: EfficientNet performance results on ImageNet

3.1.3 Preliminary preprocessing on input image

We know that we must perform preprocessing for images before feeding them into a neural network, the image must match the input size of the network.

We know that image preprocessing includes adjustments to the size, orientation, and color and some other things. We perform preprocessing to increase the image's quality in order to have a better analysis. Preprocessing aids us to eliminate unneeded distortions and noises and improve qualities which are more important for recognition or any other application that we are looking for.

All in all, for preprocessing we must perform as such as resizing, random flipping, grayscaling and etc...

3.1.4 EfficientNet Pros compared to other models

We begin by taking a look at the accuracy per number of parameters plot as given in the paper.




Figure 88: Imagenet accuracy for different models

As we can see **EfficientNet** models yield to higher accuracies with lesser

parameters compared to different models such as **ResNet**, **SENet**, etc...

Another diagram can be seen below.




Figure 89: Different models comparison

As explained in the given paper, **EfficientNet** uses compound coefficients to scale models up in a very effective manner, instead of randomly scaling width, depth or resolution scaling is done cleverly with a set of fixed scaling coefficients.

For this model, transfer learning saves our time and computational energy as do other models, but we have better accuracies do to the clever scaling method explained.

3.2 Part B : Network Implementation with transfer learning

Here we implemented the model using keras then tested it for a photo taken by myself, the codes and results are as follows.

```
1 from tensorflow.keras.utils import load_img
2 from tensorflow.keras.utils import img_to_array
3 from tensorflow.keras.applications import EfficientNetB0
4 from keras.applications.efficientnet import preprocess_input
5 from keras.applications.efficientnet import
  decode_predictions
6 modelEffnet = EfficientNetB0(weights='imagenet')
7
```

Figure 90: Defining the EfficientNetB0 model

```
1 MyImage = load_img('Teapot.jpg',target_size = (224, 224))
2 MyImage = img_to_array(MyImage)
3 MyImage = MyImage.reshape((1, MyImage.shape[0], MyImage.shape
  [1], MyImage.shape[2]))
4 MyImage = preprocess_input(MyImage)
5 ypred = modelEffnet.predict(MyImage)
6 DecodedLabels = decode_predictions(ypred)
7
```

Figure 91: Performing predictions

After this we shall print the labels with the most probability as follows.

```
1 #Most Probability #1
2 DecodedLabel1 = DecodedLabels[0][0]
3 print(f"Pred with first most probability is : {DecodedLabel1
4       [1]}\nAccuracy : {DecodedLabel1[2]*100}%")
5 #Most Probability #2
6 DecodedLabel2 = DecodedLabels[0][1]
7 print(f"Pred with second most probability is : {DecodedLabel2
8       [1]} \nAccuracy : {DecodedLabel2[2]*100}%")
9 #Most Probability #3
10 DecodedLabel3 = DecodedLabels[0][2]
11 print(f"Pred with third most probability is : {DecodedLabel3
12       [1]} \nAccuracy : {DecodedLabel3[2]*100}%")
13
14 The results:
15 Pred with first most probability is : teapot
16 Accuracy : 77.83220410346985%
17 Pred with second most probability is : coffeepot
18 Accuracy : 5.684347823262215%
19 Pred with third most probability is : water_jug
20 Accuracy : 2.3913515731692314%
```

Figure 92: Code and result for prediction printing



Figure 93: Input to the model

As we can see the predictions are close to the real thing.

3.3 Part C : Fixing a common problem

To fix the problem of input images not available in the original dataset, I have designed a function in which a threshold is observed and if the probability is less than the said threshold then the photo shall be labeled as **Others**.

The function code is depicted below.

```
1 def PhotoinModelChecker(DecodedLabels):
2     for i in range(len(DecodedLabels[0])):
3         if DecodedLabels[0][i][2] < 0.5:
4             DecodedLabels[0][i] = list(DecodedLabels[0][i])
5             DecodedLabels[0][i][1] = 'others'
6             DecodedLabels[0][i][2] = '-'
7             DecodedLabels[0][i] = tuple(DecodedLabels[0][i])
8     return DecodedLabels
9
```

Figure 94: Code and result for prediction printing

To test this I gave the following photo to the network and obtained the following results.



Figure 95: Input to the model not in the dataset

The code and results are as follows.

```
1  #Most Probability #1
2  DecodedLabel1 = DecodedLabels[0][0]
3  print(f"Pred with first most probability is : {DecodedLabel1
4  [1]}\nAccuracy : {DecodedLabel1[2]*100}%")
5  #Most Probability #2
6  DecodedLabel2 = DecodedLabels[0][1]
7  print(f"Pred with second most probability is : {DecodedLabel2
8  [1]} \nAccuracy : {DecodedLabel2[2]*100}%")
9  #Most Probability #3
10 DecodedLabel3 = DecodedLabels[0][2]
11 print(f"Pred with third most probability is : {DecodedLabel3
12 [1]} \nAccuracy : {DecodedLabel3[2]*100}%")
13
14 The results:
15 Pred with first most probability is : others
16 Accuracy : -%
17 Pred with second most probability is : others
18 Accuracy : -%
19 Pred with third most probability is : others
20 Accuracy : -%
```

Figure 96: Code and result for prediction printing

As we can see our implementation works, we shall now move on to the final part of this project.

3.4 Part 4 : Retraining the model on a new dataset

Here I used the **dogs-vs-cats** dataset which is available on Kaggle, then we perform transfer learning on this as depicted below.

```
1  ! pip install -q kaggle
2  ! mkdir ~/.kaggle
3  ! cp kaggle.json ~/.kaggle/
4  ! chmod 600 ~/.kaggle/kaggle.json
5  ! kaggle datasets download -d shaunthesheep/microsoft-
   catsvsdogs-dataset
6
```

Figure 97: Downloading the dataset from Kaggle

Two of the photos in this dataset are faulty so I have removed them.

```
1  ! unzip -q /content/microsoft-catsvsdogs-dataset.zip
2  ! rm PetImages/Cat/666.jpg
3  ! rm PetImages/Dog/11702.jpg
4  ! rm PetImages/Cat/*.db
5  ! rm PetImages/Dog/*.db
6
```

Figure 98: Deleting faulty images from dataset

Now we shall perform the train and validation split.

```
1 image_generator = ImageDataGenerator(validation_split=0.3, rescale=1/255.)
2
3 train = image_generator.flow_from_directory(batch_size=128,
4 directory='/content/PetImages',
5 shuffle=True,
6 subset="training",
7 class_mode="binary",
8 target_size=(224, 224),
9 classes = ["Cat", "Dog"])
10
11 valid = image_generator.flow_from_directory(batch_size=128,
12 directory='/content/PetImages',
13 shuffle=True,
14 subset="validation",
15 target_size=(224, 224),
16 class_mode="binary",
17 classes = ["Cat", "Dog"])
18
```

Figure 99: Train and valid split

Now we shall design the model.

```
1 base_model = EfficientNetB0(input_shape=(224,224,3), weights=
2 'imagenet', include_top=False)
3 model=Sequential()
4 model.add(base_model)
5 for layer in model.layers:
6     layer.trainable = False
7 model.add(Dropout(0.2))
8 model.add(Flatten())
9 model.add(Dense(1, kernel_initializer='he_uniform'))
10 model.add(keras.layers.Activation('sigmoid'))
11
12 model.compile(optimizer='adam', loss='binary_crossentropy',
13 metrics=['accuracy'])
14 model.summary()
```

Figure 100: The model

Now we shall train it.

```

1 137/137 [=====] - 105s 715ms/step - loss: 1.1253 - accuracy: 0.5014 - val_loss:
   0.6992 - val_accuracy: 0.5000
2 Epoch 2/20
3 137/137 [=====] - 87s 636ms/step - loss: 0.7618 - accuracy: 0.5078 - val_loss:
   0.6893 - val_accuracy: 0.5447
4 Epoch 3/20
5 137/137 [=====] - 88s 643ms/step - loss: 0.7309 - accuracy: 0.5053 - val_loss:
   0.7074 - val_accuracy: 0.5011
6 Epoch 4/20
7 137/137 [=====] - 87s 637ms/step - loss: 0.7692 - accuracy: 0.5066 - val_loss:
   0.6868 - val_accuracy: 0.5940
8 Epoch 5/20
9 137/137 [=====] - 87s 635ms/step - loss: 0.7900 - accuracy: 0.5075 - val_loss:
   1.0678 - val_accuracy: 0.5000
10 Epoch 6/20
11 137/137 [=====] - 101s 739ms/step - loss: 0.7862 - accuracy: 0.4958 - val_loss:
   0.7098 - val_accuracy: 0.5000
12 Epoch 7/20
13 137/137 [=====] - 87s 635ms/step - loss: 0.7820 - accuracy: 0.5129 - val_loss:
   0.7739 - val_accuracy: 0.5007
14 Epoch 8/20
15 137/137 [=====] - 87s 636ms/step - loss: 0.8326 - accuracy: 0.5057 - val_loss:
   0.6828 - val_accuracy: 0.5886
16 Epoch 9/20
17 137/137 [=====] - 87s 636ms/step - loss: 0.7549 - accuracy: 0.5159 - val_loss:
   0.6834 - val_accuracy: 0.5750
18 Epoch 10/20
19 137/137 [=====] - 86s 630ms/step - loss: 0.7436 - accuracy: 0.5183 - val_loss:
   0.6891 - val_accuracy: 0.5195
20 Epoch 11/20
21 137/137 [=====] - 87s 634ms/step - loss: 0.8153 - accuracy: 0.5073 - val_loss:
   0.7324 - val_accuracy: 0.5033
22 Epoch 12/20
23 137/137 [=====] - 87s 633ms/step - loss: 0.7811 - accuracy: 0.5144 - val_loss:
   0.6810 - val_accuracy: 0.5723
24 Epoch 13/20
25 137/137 [=====] - 87s 635ms/step - loss: 0.7573 - accuracy: 0.5069 - val_loss:
   0.6925 - val_accuracy: 0.5000
26 Epoch 14/20
27 137/137 [=====] - 86s 631ms/step - loss: 0.7745 - accuracy: 0.5181 - val_loss:
   0.7713 - val_accuracy: 0.5017
28 Epoch 15/20
29 137/137 [=====] - 87s 635ms/step - loss: 0.7708 - accuracy: 0.5154 - val_loss:
   0.7525 - val_accuracy: 0.5039
30 Epoch 16/20
31 137/137 [=====] - 86s 631ms/step - loss: 0.7574 - accuracy: 0.5124 - val_loss:
   0.8120 - val_accuracy: 0.5000
32 Epoch 17/20
33 137/137 [=====] - 86s 631ms/step - loss: 0.7481 - accuracy: 0.5222 - val_loss:
   0.7015 - val_accuracy: 0.5000
34 Epoch 18/20
35 137/137 [=====] - 88s 642ms/step - loss: 0.7952 - accuracy: 0.5137 - val_loss:
   0.6866 - val_accuracy: 0.5036
36 Epoch 19/20
37 137/137 [=====] - 87s 637ms/step - loss: 0.7536 - accuracy: 0.5175 - val_loss:
   0.8200 - val_accuracy: 0.5000
38 Epoch 20/20
39 137/137 [=====] - 86s 627ms/step - loss: 0.7613 - accuracy: 0.5187 - val_loss:
   0.8220 - val_accuracy: 0.5013
40

```

Figure 101: Training on dogs-vs-cats

Accuracy Plot




Figure 102: Accuracy Plot

Loss Plot




Figure 103: Loss Plot

References

- [1] [Reshad Hosseini](#), *Intelligent Systems Lecture Notes, Fall 01*
- [2] [Ali Olfat](#), *Stochastic Processes Lecture Notes, Fall 01*
- [3] [Activation functions advantages and disadvantages](#)
- [4] [Activation functions in neural networks](#)
- [5] [Batch Normalization](#)
- [6] [Batch Normalization in Keras](#)
- [7] [An introduction to pooling](#)
- [8] [Dropout in Neural Networks](#)
- [9] [Why is random forest an improvement of decision tree?](#)
- [10] [Early stopping in Keras](#)
- [11] [dogs-vs-cats dataset](#)
- [12] [EfficientNet in Keras](#)
- [13] [Image classification via fine-tuning with EfficientNet](#)
- [14] [Image Data preprocessing](#)
- [15] [Poisson Regression](#)
- [16] [Poisson Regression](#)
- [17] [Mean Squared Error vs Cross entropy loss function](#)